A Certified Core Policy Language

Bahman Sistany

Thesis submitted to the Faculty of Graduate and Postdoctoral Studies
In partial fulfillment of the requirements for the degree of Doctor of Philosophy in
Computer Science

Ottawa-Carleton Institute for Computer Science
School of Information Technology and Engineering
University of Ottawa,
Ottawa, Canada

Abstract

We present a Certified Core Policy Language (ACCPL) that can be used to express access-control rules in a minimal way. Full-blown access-control policy languages such as eXtensible Access Control Markup Language (XACML) [ANP+03] already exist however because access rules in such languages are often expressed in a declarative manner using fragments of a natural language like English, it isn't always clear what the intended behaviour of the system encoded in the access rules should be. To remedy this ambiguity, formal specification of how an access-control mechanism should behave is typically given in some sort of logic, often a subset of first order logic. To show that an access-control system actually behaves correctly with respect to its specification, proofs are needed, however the proofs that are often presented in the literature are hard or impossible to formally verify. The verification difficulty is partly due to the fact that the language used to do the proofs while mathematical in nature, utilizes intuitive justifications to derive the proofs.

In this thesis, we describe the design and implementation of ACCPL in Coq. The Coq proof assistant is used to encode and model the behaviour of ACCPL. We will use certain properties described in [TK06] to study how amenable to analysis and reasoning ACCPL policies are with respect to other access-control policy languages and how the design choices that were made contributed to this ease of reasoning. In particular decidability criteria for ACCPL are expressed and proofs are developed in the Coq proof assistant.

Table of Contents

Listings

Chapter 1

Introduction

1.1 Access-Control Models

Authorization refers to the process of rendering a decision to allow access to a resource or asset of interest. By the same token all unauthorized access requests to resources must be controlled and ultimately denied, hence the term "access-control".

Various access-control models exist and below we give a short summary of some of the most important ones leading up to our own core language, ACCPL.

1.1.1 Access Control Lists

Access Control List (ACL) is perhaps the oldest and the most basic access-control model. A list of subjects along with their rights are kept per resource or object of interest. Every time a subject makes an access request on an object, the ACL of the object is consulted and access is either granted or denied to the requester based on whether the requester is listed in the ACL and has the correct set of rights.

1.1.2 Capabilities

Capabilities based access-control works based on a list of objects and associated rights. The list of objects and the associated rights comprise an "unforgeable" ticket that a reference monitor checks to allow access (or not). Capabilities based systems don't need to authenticate users as ACL based systems do.

1.1.3 RBAC

In Role-based Access Control (RBAC) systems a requester's role determines whether access is granted or denied. In this model, users belong to roles and rights are associated with roles so no direct association between users and rights exist. Roles are meant to group users

and add flexibility when assigning rights. RBAC systems naturally solve the problem of assigning ACLs for a large group of users and manage the administration cost of changing users' rights in ACL based systems.

1.1.4 ABAC

Despite many advantages of RBAC systems, some disadvantages also exist. Often a role needs to be decomposed into sub-roles based on the type of resource to be administered, and perhaps also based on the location that the resource serves [NIS09]. Basically RBAC suffers from a lack of a sub-typing mechanism whereby individual members of a group/role may be differentiated and access is granted or denied based on a more granular set of attributes. Attribute-based Access Control (ABAC) model was proposed to fulfill this granularity requirement. In ABAC access control decisions are made based on a set of attributes, associated with the subject making the request, the environment, and/or the resource itself [NIS09].

1.1.5 PBAC

In order to harmonize access control in large environments with many subjects and objects and disparate attributes, Policy-based Access Control (PBAC) model has been proposed [NIS09]. PBAC allows for a more uniform access-control model across the system. PBAC systems help create and enforce policies that define who should have access to what resources, and under what circumstances [NIS09]. There is also a need for large organizations to put in place mechanisms such that access-control rules can be easily audited. This calls for a data-driven approach to access-control where the data, in this case the access-control rules, are available to read and analyze. A data-driven approach like PBAC additionally helps with modularity of the system as changes in access control rules will have almost no impact to the underlying system the rules are meant to protect.

1.2 A Core Policy Language for PBAC Systems

Does there already exist a suitable policy language that can be used for expressing general access-control expressions? How about XACML, Open Digital Rights Language (ODRL) [Ian02] or eXtensible rights Markup Language (XrML) [WLD$^+$02]?

Currently the most popular Rights Expression Languages (REL)s are the XrML, and the ODRL. Both of these languages are XML based and are considered declarative languages. XrML has been selected to be the REL for *MPEG-21* which is an ISO standard for multimedia applications. ODRL is also a standards based REL which has been accepted as part of the W3C community with the mandate of standardizing how rights and policies, related to the usage of digital content on the Open Web Platform, *OWP* [Wik15b], are expressed. ODRL 2.0 supports expression of rights and also privacy rules for social media

while ODRL 1.0 was only dealing with the mobile ecosystem – ODRL 1.0 was adopted by the Open Mobile Alliance, *OMA* in 2000.

RELs, or more precisely Digital Rights Expression Languages (DREL)s when dealing with digital assets deal with the "rights definition" aspect of the Digital Rights Management (DRM) ecosystem. A DREL, allows the expression and definition of digital asset usage rights such that other areas of the DRM ecosystem, namely the enforcement mechanism and the usage tracking components can function correctly.

As popular as both XrML and ODRL are, their adoption and usage is still somewhat limited in practice. Both Apple and Microsoft for example have defined their own lightweight RELs [JHM06] in *Fair Play* [Wik15a] and in *PlayReady* [Wik15c]. The authors of [JHM06] argue that both these RELs (XrML and ODRL) and others are simply too complex to be used effectively (for expressing rights) since they also try to cover much of the the enforcement and tracking aspects of DRMs.

DRM refers to the digital management of rights associated with the access or usage of digital assets. There are various aspects of rights management however. According to the authors of the white paper "A digital rights management ecosystem model for the education community,"([CPR04]) digital rights management systems cover the following four areas: 1) defining rights 2) distributing/acquiring rights 3) enforcing rights and 4) tracking usage.

XACML is a high-level and platform independent access control system that is also XML based. XACML is an OASIS standard which defines a language for the definition of policies and access requests, and a workflow to achieve policy enforcement [MPT12]. According to Masi et al [MPT12] designing XACML access control policies is difficult and error prone. Furthermore XACML comes without a formal semantics as do both XrML and ODRL. The XACML standard is written in prose and contains quite a number of loose points that may give rise to different interpretations and lead to different implementation choices [MPT12].

XACML, ODRL and XrML are all PBAC based languages where ODRL and XrML differ from XACML by their focus on digital assets protection and in general DRM, hence the term REL. All three are full-blown and custom languages that have one thing in common; they suffer from a lack of formal semantics. Additionally all of these languages cover much more than policy expressions leading to access decisions; they also address enforcement of the policies (ODRL and XrML specifically and DRM in general distinguish themselves from general access-control languages by additionally addressing enforcement of policies beyond where the policies were generated). A third reason that made these custom languages unsuitable as a core policy language was the fact that they are limited in terms of what can be built on top of them; for example expressing hierarchical role-based access-control in XACML requires a fairly complex encoding [TK06].

Rights expressions in DRELs and specifically ODRL are used to arbitrate access to assets under conditions. The main construct in ODRL is the *agreement* which specifies users, asset(s) and policies whereby controls on users' access to the assets are described. This is very similar to how access control conditions are expressed in access control policy languages such as XACML [ANP⁺03] and SELinux [SVLS02]. In fact several authors have

3

worked on interoperability between RELs and access control policy languages, specifically between ODRL and XACML, [PRD05, MRD09] and also on translation from high level policies of XACML to low-level and fine grained policies of SELinux [ASLZ08].

A policy language that was based on some sort of logic with formal semantics and also one that was minimal and extendible was clearly needed. We started by looking at Lithium [HW08] and subsequently Pucella and Weissman's subset of ODRL [PW06] as potential core languages. Lithium uses DRM as the main application of its access-control system whereas Pucella and Weissman's primary goal is to define formal semantics for ODRL whose main application in turn is DRM. We will use Pucella and Weissman's subset of ODRL as the basis for ACCPL and in doing so treat digital rights as our main access-control application without loss of generality with respect to other applications, with the final goal of performing formal verification on policies written in ACCPL.

1.3   Formal Semantics for PBAC Languages

Formal methods help ensure that a system behaves correctly with respect to a specification of its desired behavior [Pie02]. This specification of the desired behavior is what's referred to as *semantics* of the system. Using formal methods requires defining precise and formal semantics, without which analysis and reasoning about properties of the system in question would become impossible. For example, an issue with the current batch of RELs are due to their semantics being expressed in a natural language (e.g. English) which by necessity results in ambiguous and open to interpretation behavior.

To formalize the semantics of PBAC languages several approaches have been attempted by various authors. Most are logic based [HW08, PW06] while others are based on finite-automata [HKS04], operational semantics based interpreters [SS09] and web ontology (from the Knowledge Representation Field) [KG10].

1.4   Logic Based Semantics

Formal logic can represent the statements and facts we express in a natural language like English. Propositional logic is expressive enough to express simple facts as propositions and uses connectives to allow for the negation, conjunction and disjunction of the facts. In addition simple facts can be expressed conditionally using the implication connective. Propositional logic however is not expressive enough to express policies of the kind used in languages like ODRL and XrML. For example, a simple policy expressed in English like "All who pay 5 dollars can watch the movie Toy Story" cannot be expressed in propositional logic because the concept of variables doesn't exist in propositional logic.

A richer logic such as "Predicate Logic", also called "First Order Logic" (*FOL*), is more suitable and has the expressive power to represent policies written in English. Moreover, FOL can be used to capture the meaning of policies in an unambiguous way.

Halpern and Weissman [HW08] propose a fragment of FOL to represent and reason about policies. The fragment of FOL they arrive at is called *Lithium* which is decidable and allows for efficiently answering interesting queries. Lithium restricts policies to be written based on the concept of "bipolarity" which disallows by construction policies that both permit and deny an action on an object. Pucella and Weissman [PW06] specify a predicate logic based language that represents a subset of ODRL.

1.5 Specific Problem

Policy languages and the agreements written in those languages are meant to implement specific goals such as limiting access to specific assets. The tension in designing a policy language is usually between how to make the language expressive enough, such that the design goals for the policy language may be expressed, and how to make the policies verifiable with respect to the stated goals.

As stated earlier, an important part of fulfilling the verifiability goal is to have formal semantics defined for policy languages. For ODRL, authors of [PW06] have defined a formal semantics based on which they declare and prove a number of important theorems (their main focus is on stating and proving algorithm complexity results). However as with many paper-proofs, the language used to do the proofs while mathematical in nature, uses a lot of intuitive justifications to show the proofs. As such these proofs are difficult to verify or more importantly to "derive". Furthermore the proofs can not be used directly to render a decision on a sample policy (e.g. whether to allow or deny access to an asset). Of course one may (carefully) construct a program based on these proofs for practical purposes but certifying such programs correct presents additional verification challenges, even assuming the original proofs were in fact correct.

1.6 Contributions

In this thesis we have built a language called ACCPL based on ODRL and starting with definitions in [PW06]. The ACCPL framework has been encoded in *Coq* [BC04] which is both a programming language and a proof-assistant. We have declared and proved decidability results for ACCPL in Coq which will allow us to extract programs from the proofs; the executable programs can be used on specific policies to render a specific decision such as "a permission has been granted".

We originally started with a specific subset of [PW06] so that we could concentrate on what we believed to be the essence or core of the language. For example, we started with only one of three different kinds of "facts" (that may affect the permit/deny type decisions). We also had to change some of the language productions to allow for Coq's requirement for clearly terminating recursion. However we maintained the central semantic definitions including "Closed World Assumptions" [PW06] where the semantics only specify explicitly Permitted and notPermitted answers so the semantics as stated by Pucella and Weissman [PW06] are not complete and therefore not decidable. We have made

major modifications to the semantics of Pucella and Weissman's language to make the new ACCPL language decidable.

Our decidability results subsume an important sub-category, namely inconsistency or conflict-detection in policy expressions or rules. Authors of [SM12] [SMF13] describe and implement in Coq a conflict detection algorithm for detecting conflicts in XACML access control rules. XACML is an expressive and at the same time complex policy language which makes conflict detection a difficult task. Authors of [SM12] then prove the conflict detection algorithm correct (or certified) by developing a formal proof in Coq. The proof is rather complex and involves a large number of cases, including many corner cases that were difficult to get right [SM12]. For ACCPL we have formally proven that conflicts are not possible.

ACCPL being a core policy language with certified semantics could be used to implement various policy languages and reason about interoperability between those languages [PRD05, MRD09]. In this manner our Coq based language ACCPL can be viewed as *abstract syntax*, complete with defined semantics, that can be used for implementing various policy languages with more concrete syntax (e.g. W3C's ODRL and SELinux).

For access to the Coq source code for ACCPL, please refer to the GitHub repository at https://github.com/bsistany/yacpl.git.

Chapter 2

ACCPL Syntax

2.1 Introduction

We follow the style of [PW06] by using abstract syntax to express policy statements in ACCPL. Abstract syntax is a more compact representation than XML which is what all the XML-based policy languages such ODRL use. Furthermore abstract syntax simplifies specifying the semantics as we shall see later. As an example the agreement "the asset TheReport may be printed a total of 5 times by either Alice or Bob and twice more by Alice only" written in ODRL's XML encoding is illustrated in Listing 2.1 [PW06].

Listing 2.1: Agreement for Alice and Bob in XML

```
<agreement>
 <asset> <context> <uid> The Report </uid> </context> </asset>
 <permission>
   <print>
    <constraint> <count> 5 </count> </constraint>
   </print>
   <print>
    <constraint> <name> Alice </name> </constraint>
    <constraint> <count> 2 </count> </constraint>
   </print>
 </permission>
 <party> <context> <name> Alice </name> </context> </party>
 <party> <context> <name> Bob </name> </context> </party>
</agreement>
```

The agreement in Listing 2.1 is shown in 2.2 using the syntax from [PW06].

Listing 2.2: Agreement for Mary Smith as BNF (as used in [PW06])

```
agreement
 for Alice and Bob
 about The Report
```

```
  with True -> and[count[5] => print], and[Alice, count[2] =>
    print].
```

In the following we will cover the *abstract syntax* of ACCPL that we later implement in Coq using constructs such as *Inductive Types* and Definitions.

2.2   Environmental Facts

In ACCPL, agreements and facts (i.e. environments) will refer to a count of how many times each policy should be and has been used to justify an action. This is the only fact that ACCPL will cover although we conjecture adding other facts and the machinery to support those facts should not change verification goals and results so far of ACCPL.

In ACCPL a *prerequisite* is either *true*, a *constraint*, the negative of a *constraint* or a conjunction of *prerequisite*s. *true* is the prerequisite that always holds. Constraints are facts that are outside of control of users. For example, there is nothing *Alice* can do to satisfy the constraint "user must be Bob".

We will describe ACCPL in a *BNF* grammar that looks more like Pucella and Weissman's subset grammar [PW06]. BNF style grammars are more abstract as they only give suggestions about the surface syntax of expressions without getting into lexical analysis and parsing related aspects such as precedence order of operators [PCG+11]. The Coq version in contrast is the encodings of the BNF definitions which is tyically used for building compilers and interpreters. We will present both the BNF definition and its Coq encoding for each construct of ACCPL.

2.3   Productions

The top level ACCPL production is the *agreement*. An agreement expresses what actions a set of subjects may perform on an object and under what conditions. Syntactically an agreement is composed of a set of subjects/users called a *principal* or *prin*, an *asset* and a *policySet*.

Listing 2.3: agreement

```
<agreement> ::=
    'agreement' 'for' <prin> 'about' <asset> 'with' <
      policySet>
```

Principals or prins are composed of *subjects* which are specified based on the application e.g. Alice, Bob, etc.

Listing 2.4: prin

```
<prin> ::=  { <subject_1>, ..., <subject_m> }
```

```
<subject> ::= N
```

Assets are also application specific but similar to subjects we will continue using specific ones for the DRM application (taken from [PW06]). *ebook*, *The Report* and *latestJingle* are examples of specific subjects we will be using throughout. Syntactically an asset is represented as a natural number (N). Similarly for subjects.

Listing 2.6: asset

```
<asset> ::= N
```

Agreements express who may perform an action on an asset. They include a set of subjects (i.e. a *prin*), an asset and a policy set. A policy set is a primitive policy set implying non-nested policy sets. Note that we could define various combining operators for policy sets such as conjunctions and disjunctions but we keep ACCPL's policy sets limited to the primitive kind but use a policy combining operator when it comes to dealing with policies later on. Each primitive policy set specifies a *prerequisite* and a *policy*. In general if the prerequisite "holds" the policy is taken into consideration. Otherwise the policy will not be looked at. Some primitive policy sets are specified as *inclusive* as opposed to others that are explicitly specified as *exclusive*. The *Primitive Exclusive Policy Set*s are exclusive to agreement's users in that only those users may perform the actions specified in the policy set. The implication is that all other users who are not specified in the agreement's principal (prin) are forbidden from performing the specified actions, no matter whether the *prerequisite* holds or not. Not surprisingly we also define *Primitive Inclusive Policy Set*s that don't enforce any exclusivity to the agreement's users.

Listing 2.7: policySet

```
<policySet> ::=
    <primPolicySet>                         ; primitive policy set
```

Listing 2.8: primPolicySet

```
<primPolicySet> ::=
    <primInclusivePolicySet>        ; primitive inclusive policy set
    <primExclusivePolicySet>        ; primitive exclusive policy set
```

Listing 2.9: primInclusivePolicySet

```
<primInclusivePolicySet> ::=
    <preRequisite> → <policy>       ; primitive inclusive policy set
```

Listing 2.10: primExclusivePolicySet

```
<primExclusivePolicySet> ::=
    <preRequisite> ↦ <policy>       ; primitive exclusive policy set
```

A primitive policy specifies an action to be performed on an asset, depending of whether the policy's prerequisite holds or not. If the prerequisite holds the agreement's user is permitted to perform the action on the agreement's asset; otherwise permission is denied. Pucella and Weissman's subset of ODRL [PW06] specify a unique identifier for each policy to help the translation (from agreements to formulas). ACCPL has maintained the identifier for future work and we include it here in our definition of the policy construct however as far as the proofs are concerned the policy identifier could be removed without a loss to the obtained results. Primitive policies could also be grouped together using the conjunction combining operator.

Listing 2.11: policy

```
<policy> ::=
 'and'[ <primPolicy_1>, ...,
                  <primPolicy_m> ]    ; conjunction
```

Listing 2.12: primPolicy

```
<primPolicy> ::=
   <preRequisite> ⇒_<policyId> <act>    ; primitive policy
```

An *Action* (*act*) is represented as a natural number. Similar to assets and subjects, actions are application specific. Some example actions taken from [PW06] are *Display* and *Print*.

Listing 2.13: act

```
<act> ::= N
```

A *Policy Id* (*policyId*) is a unique identifier specified as (increasing) positive integers.

Listing 2.14: policyId

```
<policyId> ::= N
```

In ACCPL a *prerequisite* is either true or it is a *constraint*. The *true* prerequisite always holds. A constraint is an intrinsic part of a policy and cannot be influenced by agreement's user. Minimum height requirements for popular attractions and rides are examples of what we would consider a constraint. *NotCons* is a negation of a constraint. Finally the set of prerequisites are closed under conjunction operator (*AndPrqs*).

Listing 2.15: preRequisite

```
<preRequisite> ::=
        'True'                        ; always true
        <constraint>                  ; constraint
        'not' [ <constraint> ]        ; suspending constraint
        'and'[ <preRequisite_1>, ..., <preRequisite_m> ]
            ; conjunction
```

Constraints are either *Principal, Count* or *CountByPrin.* Principal constraints basically require matching to specified prins. For example, the user being Alice is a Principal constraint. A count constraint refers to a set of policies P and specifies the number of times the user of an agreement has invoked the policies in P to justify her actions. If the count constraint is part of a policy then the set P is composed of the single policy. In the case that the count constraint is part of a policy set, the set P is the set of policies specified in the policy set.

Listing 2.16: constraint

```
<constraint> ::=
        <prin>                         ; principal
        'Count' [N]                    ; number of executions
        <prin> ('Count' [N])           ; number of executions by prin
```

Chapter 3

ACCPL Syntax In Coq

3.1 Introduction to Coq

Coq is known first and foremost as a proof-assistant. The underlying formal language that Coq uses is a much more expressive version of typed lambda calculus called Calculus of (Co)Inductive Constructions (CIC) where proofs and programs can both be represented. For example, CIC adds polymorphism (terms depending on types), type operators (types depending on types) and dependent types (types depending on terms).

Specifications of programs in Coq may be expressed using the specification language *Gallina* [Hue92]. Coq is then used to develop proofs to show that a program's run-time behaviour satisfies its specification. Such programs are called *certified* because they are formally verified and confirmed to conform to their specifications [BC04].

Assertions or propositions are statements about values in Coq such as $3 < 8$ or $8 < 3$ that may be true, false or even be only conjectures. To verify that a proposition is true a proof needs to constructed. While paper-proofs use a combination of mathematics and natural language to describe their proofs, Coq provides a formal (and therefore unambiguous) language that is based on proof-theory to develop proofs in. Verification of complex proofs is possible because one can verify the intermediate proofs or sub-goals in steps, each step being derived from the previous by following precise derivation rules. The Coq proof engine solves successive goals by using predefined *tactics*. Coq tactics are commands to manipulate the local context and to decompose a goal into simpler goals or sub-goals [BC04].

3.2 ACCPL Syntax

ACCPL productions were presented as high level abstract syntax in 2.3. Below we present the corresponding encodings in Coq.

An agreement is a new inductive type in Coq by the same name. The constructor *Agreement* takes a *prin*, an *asset* and a *policySet*. *prin* is defined to be a non empty list of *subject*s (see listing 3.1).

Types *asset*, *subject*, *act* and *policyId* are simply defined as *nat* which is the datatype of natural numbers defined in coq's library module *Coq.Init.Datatypes* (*nat* is itself an inductive datatype). We use Coq constants to refer to specific objects of each type. For example, the subject 'Alice' is defined as *DefinitionAlice* : *subject* := 101. and the act 'Play' as *DefinitionPlay* : *act* := 301.. For each "nat" type in ACCPL we have also used constants that play the role of "Null" objects (see "Null Object Pattern" [MRB98]), for example *NullSubject*. This is needed partly because of the way ACCPL language elements are defined which corresponds to the need to use *nonemptylist* exclusively even though at intermediate stages during the various algorithms Coq's *list* is a better fit because it allows empty lists.

Next we define the *policySet* datatype. Note the close/one-to-one mapping to its counterpart in listing 2.7. A *policySet* is constructed only one way: by calling the *PPS* constructor which takes a *primPolicySet* as input. There are two ways a *primPolicySet* can be constructed (see listing 2.8) corresponding to two constructors: *PIPS* and *PEPS*.

PIPS takes a *primInclusivePolicySet* as input while *PEPS* takes a *primExclusivePolicySet*. Both *primInclusivePolicySet* and *primExclusivePolicySet* types are constructed by taking a *preRequisite* and a *policy* as parameters (see listing 2.9 and 2.10).

A *policy* is defined as a datatype with the constructors *Policy* which takes a *nonemptylist* (see listing 3.3) of *primPolicy*s (see listing 2.11). *primPolicy* is constructed by calling *PrimitivePolicy* which takes a *preRequisite*, a *policyId* and an action *act* (see listing 2.12). Ignoring the *policyId* for a moment, a primitive policy consists of a prerequisite and an action. If the prerequisite holds the action is allowed to be performed on the asset.

The data type *nonemptylist* reflects the definition of "policy conjunction" in listing listing 2.11 in chapter 2 (see definition of *nonemptylist* in listing 3.3). Essentially *nonemptylist* represents a list data structure that has at least one element and it is defined as a new *polymorphic* inductive type in its own section.

Listing 3.1: ACCPL: Coq Version of Agreement

```coq
Inductive agreement : Set :=
  | Agreement : prin → asset → policySet → agreement.

Definition prin := nonemptylist subject.

Definition asset := nat.

Definition subject := nat.

Definition act := nat.

Definition policyId := nat.

Inductive policySet : Set :=
  | PPS : primPolicySet → policySet.

Inductive primPolicySet : Set :=
  | PIPS : primInclusivePolicySet → primPolicySet
  | PEPS : primExclusivePolicySet → primPolicySet.

Inductive primInclusivePolicySet : Set :=
  | PrimitiveInclusivePolicySet : preRequisite → policy →
    primInclusivePolicySet.

Inductive primExclusivePolicySet : Set :=
  | PrimitiveExclusivePolicySet : preRequisite → policy →
    primExclusivePolicySet.

Inductive policy : Set :=
  | Policy : nonemptylist primPolicy → policy.

Inductive primPolicy : Set :=
  | PrimitivePolicy : preRequisite → policyId → act → primPolicy.
```

Listing 3.2: ACPL: Coq Version of Agreement

```coq
Inductive agreement : Set :=
  | Agreement : prin → asset → policySet → agreement.

Definition prin := nonemptylist subject.

Definition asset := nat.

Definition subject := nat.

Definition act := nat.

Definition policyId := nat.

Inductive primPolicy : Set :=
  | PrimitivePolicy : preRequisite → policyId → act → primPolicy.

Inductive policy : Set :=
  | Policy : nonemptylist primPolicy → policy.

Inductive primInclusivePolicySet : Set :=
  | PrimitiveInclusivePolicySet : preRequisite → policy →
    primInclusivePolicySet.

Inductive primExclusivePolicySet : Set :=
  | PrimitiveExclusivePolicySet : preRequisite → policy →
    primExclusivePolicySet.

Inductive primPolicySet : Set :=
  | PIPS : primInclusivePolicySet → primPolicySet
  | PEPS : primExclusivePolicySet → primPolicySet.

Inductive policySet : Set :=
  | PPS : primPolicySet → policySet.
```

Listing 3.3: nonemptylist

```
Section nonemptylist.

Variable X : Set.

Inductive nonemptylist : Set :=
  | Single : X → nonemptylist
  | NewList : X → nonemptylist → nonemptylist.

End nonemptylist.
```

In listing 3.5 *preRequisite* is defined as a new datatype with constructors *TruePrq*, *Constraint*, *NotCons* and *AndPrqs* (see listing 2.15 for the abstract syntax equivalent).

TruePrq represents the always true prerequisite. The *Constraint* prerequisite is defined as the type *constraint* so its description is deferred here. Intuitively a constraint is a prerequisite to be satisfied that is outside the control of the user(s). For example, the constraint of being 'Alice' if you are 'Bob' (or 'Alice' for that matter). The constructor *NotCons* is defined the same way the *Constraint"constructor* is. This constructor is defined as the type *constraint* and it is meant to represent the negation of a *constraint* as we shall see in the translation (see listing 4.24). The remaining constructor *AndPrqs* takes as parameters non empty lists of prerequisites. This constructor represents the conjunction combining operator.

Finally a *constraint* (see listing 2.16 for the abstract syntax equivalent) is defined as a new datatype with constructors *Principal*, *Count* and *CountByPrin*.

Principal constraint takes a *prin* to match. For example, the constraint of the user being Bob would be represented as "Principal constraint". The *Count* constructor takes a *nat* which represents the number of times the user of an agreement has invoked the corresponding policies to justify her actions. If the count constraint is part of a policy then the corresponding policies is basically the single policy, whereas in the case that the count constraint is part of a policy set, the corresponding policies would be the set of those policies specified in the policy set. The *CountByPrin* is similar to *Count* but it takes an additional *prin* parameter. In this case the subjects specified in the *prin* parameter override the agreements' user(s).

Listing 3.4: preRequisite

```
Inductive preRequisite : Set :=
  | TruePrq : preRequisite
  | Constraint : constraint → preRequisite
  | NotCons : constraint → preRequisite
  | AndPrqs : nonemptylist preRequisite → preRequisite.



Inductive constraint : Set :=
  | Principal : prin → constraint
  | Count : nat → constraint
  | CountByPrin : prin → nat → constraint.
```

We started with the encoding of the agreement for Mary Smith in chapter 2 but we deferred the definition of the Coq constructs used to define that agreement. All the definitions needed to encode the agreement for Mary Smith except for the prerequisite *PrePay* (which is not part of ACCPL) have now been defined. More agreement examples will be given in Chapter 7.

Listing 3.5: preRequisite

```
Definition p1A1:primPolicySet :=
 PIPS (PrimitiveInclusivePolicySet
   TruePrq
   (Policy (Single (PrimitivePolicy (Constraint (Count 5)) id1 Print)))).


Definition p2A1prq1:preRequisite := (Constraint (Principal (Single Alice))).
Definition p2A1prq2:preRequisite := (Constraint (Count 2)).

Definition p2A1:primPolicySet :=
 PIPS (PrimitiveInclusivePolicySet
   TruePrq
   (Policy
     (Single
       (PrimitivePolicy (AndPrqs (NewList p2A1prq1 (Single p2A1prq2))) id2 Print)))).



Definition A1 := Agreement (NewList Alice (Single Bob)) TheReport (PPS p1A1).
```

Chapter 4

ACCPL Semantics In Coq

4.1   Introduction

The translation functions plus the auxiliary types and infrastructure which implement the semantics for ACCPL have been encoded in Coq. Translation functions build Coq terms of type *Prop* for the most part. Well-formed propositions (or *Prop*s) are assertions one can express about values such as mathematical objects or even programs (e.g. $3 < 8$) in Coq.

Whether a permission is granted or denied depends on the agreements in question but also on the facts recorded in the environment. For ACCPL those facts revolve around the number of times a policy has been used to justify an action (see section 2.2 for more details on ACCPL). We encode this information in an *environment* which is a conjunction of equalities of the form $count(s, policyId) = n$.

The Coq version of the count equality is a new inductive type called *count_equality*. An environment is defined to be a non-empty list of *count_equality* objects (see listing 4.1). Function *make_count_equality* in listing 4.1 is simply a convenience function that builds *count_equality*s. For an example of how environments are created see listing 4.2.

Listing 4.1: Environments and Counts

```
Inductive count_equality : Set :=
  | CountEquality : subject → policyId → nat → count_equality.

Definition make_count_equality
  (s:subject)(id:policyId)(n:nat): count_equality :=
  CountEquality s id n.

Inductive environment : Set :=
  | SingleEnv : count_equality → environment
  | ConsEnv :  count_equality →  environment → environment.
```

Listing 4.2: Defining Environments

```
Definition e1 : environment :=
  (SingleEnv (make_count_equality Alice id1 8)).
```

We also define a *getCount* function (see listing 4.3) that given a pair consisting of a subject and policy id, looks for a corresponding count in the environment. *getCount* assumes the given environment is consistent (meaning it won't return two different counts for the same pair of subject and policy id), so it returns the first matched *count* it sees for a $(subject, id)$ pair. If a *count* for a $(subject, id)$ pair is not found it returns 0.

Listing 4.3: getCount Function

```
Fixpoint getCount
  (e:environment)(s:subject)(id: policyId): nat :=
 match e with
 | SingleEnv f  ⇒
     match f with
     | CountEquality s1 id1 n1 ⇒
         if (beq_nat s s1)
         then if (beq_nat id id1) then n1 else 0
         else 0
     end
 | ConsEnv f rest ⇒
     match f with
     | CountEquality s1 id1 n1 ⇒
         if (beq_nat s s1)
         then if (beq_nat id id1) then n1 else (getCount rest s id)
         else (getCount rest s id)
     end
 end.
```

4.2   Decision Procedures and the sumbool type

sumbool is a boolean type defined in the Coq standard library module $Coq.Init.Specif$. sumbool is equipped with the justification of its value, which really helps with proofs. sumbool A B, which is written {A}+{B}, is the informative disjunction "A or B", where A and B are logical propositions. Its extraction is isomorphic to the type of booleans. A boolean is either true or false, and this is decidable [dt04]. See listing 4.4 for the definition of *sumbool* and *sumbool_of_bool*. With sumbools similar to bool and other 2-constructor inductive types, one can use the "if then else" construct to select the desired case when doing proofs which makes for much more readability of the given proofs.

Listing 4.4: sumbool type

```
Inductive sumbool (A B:Prop) : Set :=
  | left : A → {A} + {B}
  | right : B → {A} + {B}
 where "{ A } + { B }" := (sumbool A B) : type_scope.


Definition sumbool_of_bool : ∀ b:bool, {b = true} + {b = false}.
```

We have used the sumbool type to declare and prove decision procedures that we have subsequently used in the semantics and also in the proofs. In the following we will list only the type declaration of each decision procedure and the accompanying functions. The individual decision procedures were used as a building block to make the final decidability proofs for ACCPL possible. We will refer to these listings later when we discuss the translations.

Listing 4.5: Decision Procedures: eq_nat_dec [dt04]

```
Theorem eq_nat_dec : ∀ n m, {n = m} + {n <> m}.
```

Listing 4.6: Decision Procedures: is_subject_in_prin

```
Fixpoint is_subject_in_prin (s:subject)(p:prin): Prop :=
 match p with
 | Single s' ⇒ s=s'
 | NewList s' rest ⇒ s=s' \/ (is_subject_in_prin s rest)
 end.

Theorem subject_in_prin_dec :
   ∀ (a:subject) (l:prin), {is_subject_in_prin a l} + {~ is_subject_in_prin a l}.
```

Listing 4.7: Decision Procedures: trans_prin_dec

```
Theorem trans_prin_dec :
  ∀ (x:subject)(p: prin), {trans_prin x p} + {~trans_prin x p}.
```

Listing 4.8: Decision Procedures: trans_count_dec

```
Theorem trans_count_dec :
   ∀ (n:nat), {trans_count n} + {~trans_count n}.
```

Listing 4.9: Decision Procedures: trans_constraint_dec

```
Theorem trans_constraint_dec :
  ∀ (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)(prin_u:
   prin),
    {trans_constraint e x const IDs prin_u} + {~trans_constraint e x const IDs
   prin_u}.
```

```
Theorem double_neg_constraint:
  ∀ (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)(prin_u:
   prin),
    (trans_constraint e x const IDs prin_u) → ˜˜(trans_constraint e x const IDs
   prin_u).


Theorem trans_negation_constraint_dec :
  ∀ (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)(prin_u:
   prin),
    {˜trans_constraint e x const IDs prin_u} + {˜˜trans_constraint e x const IDs
   prin_u}.
```

```
Theorem trans_notCons_dec :
  ∀ (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)(prin_u:
   prin),
    {trans_notCons e x const IDs prin_u} + {˜ trans_notCons e x const IDs prin_u}.
```

```
Theorem trans_preRequisite_dec :
  ∀ (e:environment)(x:subject)(prq:preRequisite)(IDs:nonemptylist policyId)(prin_u:
   prin),
    {trans_preRequisite e x prq IDs prin_u} + {˜ trans_preRequisite e x prq IDs
   prin_u}.
```

4.3   answer and result types

The policy translation functions ultimately will return one of the following answers: Permitted, NotPermitted or Unregulated. Permitted answer signifies that the access request has been granted. The NotPermitted answer is when the access granted is denied. And finally the Unregulated answer is for all the times when neither Permitted or NotPermitted answers are applicable. The *answer* type in the Coq listing 4.13 implements the answer concept. We wrap *answer*s in a *result* type and add some context. Intuitively a result will tell us whether a subject may perform an action on an asset or not, or else that the request is unregulated.

Listing 4.13: Decision Procedures:

```
Inductive answer : Set :=
  | Permitted : answer
  | Unregulated : answer
  | NotPermitted : answer.

Inductive result : Set :=
  | Result : answer → subject → act → asset → result.
```

4.4 Translations

Translation of the top level *agreement* element proceeds by case analysis on the structure of the agreement. However an agreement can only be built one way; by calling the constructor *Agreement*. The translation proceeds by calling the translation function for the corresponding *policySet* namely the parameter to *Agreement* called *ps*. See listing 4.14. The formal argument *e* of type *environment* is passed along many translation function but will only eventually be used to get the count information from the *getCount* function. Other formal arguments that are passed are the agreement itself to be translated, plus the subject, action and asset coming from a "query" or request for access: *action_from_query*, *subject_from_query* and *asset_from_query*. Notice that the translation function for agreements returns a nonempty list of *result*s.

4.4.1 Policy Combinators

The nonempty list of *result*s returned by the agreement translation function will have a result per *primPolicy* in the agreement. As we will see later when we discuss the proofs, we have proven that results containing a Permitted answer and a NotPermitted answer are mutually exclusive. Therefore existence of a single (or more) Permitted type result in the nonempty list makes the set "Permitted" whereas the existence of a single (or more) NotPermitted type result makes the whole set "NotPermitted". When no Permitted or NotPermitted is seen in the set, we only have Unregulated results which would make the whole set "Unregulated".

Listing 4.14: Translation of Agreement

```
Definition trans_agreement
  (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset) : nonemptylist result :=

  match ag with
  | Agreement prin_u a ps ⇒
      (trans_ps e action_from_query subject_from_query asset_from_query ps prin_u a)
  end.
```

Translation of a *policySet* (called *trans_ps* in listing 4.20), takes as input e, the environment, the subject, action and asset coming from a query: *action_from_query*, *subject_from_query* and *asset_from_query*, ps, the policy set, *prin_u*, the agreement's user, and a, the asset mentioned in the agreement.

The translation starts with checking that there is a proof for the equality of asset_-from_query and a (the asset from the agreement). If so, the translation function recurses on the composing *primPolicySet*. Otherwise, a result of Unregulated is immediately is returned. Intuitively we are just explicitly saying that the semantics won't apply to queries about assets that are not mentioned in the agreement.

A *primPolicySet* was either constructed by a *PIPS* or by a *PEPS* which distinguish between *primInclusivePolicySet* and *primExclusivePolicySet* subtypes. The translation functions *trans_policy_PIPS* and *trans_policy_PEPS* are called in turn as the case may be (see listing 4.19).

trans_policy_PIPS translation function for a *primInclusivePolicySet* checks whether there is a proof for the subject in question, x, being in *prin_u*. If so there is a check for whether the *preRequisite* from the policySet holds or not. If so the translation function *trans_policy_positive*(4.17) is called. In all other cases *trans_policy_unregulated*(4.15) is called.

trans_policy_PEPS translation function for a *primExclusivePolicySet* checks whether there is a proof for the subject in question, x, being in *prin_u*. If not, the translation function *trans_policy_negative* (4.16) is called since the policySet is exclusive to agreement's users we expect that all subjects not in the *prin_u* are denied access. In case the check for x, being in *prin_u* succeeds, there is a check for whether the *preRequisite* from the policySet holds or not. If so the translation function *trans_policy_positive* is called, else *trans_policy_unregulated* is returned.

trans_policy_positive, *trans_policy_negative* and *trans_policy_unregulated* translation functions start by recursing on a list of *primPolicy*s where the results of processing each *primPolicy* is appended to the rest of the set (see high level pseudocode in listing 4.18) for the common pattern).

trans_policy_positive acting on a single *primPolicy* (see process_single_pp_trans_-policy_positive in listing 4.17), first checks whether the *preRequisite* from the policy holds or not. If so the translation function checks that there is a proof for the equality of action_-from_query and action (the action from the agreement). If so, the translation function finally returns a result of Permitted. In all other cases, a result of Unregulated is returned. Intuitively we are just explicitly saying that the semantics won't apply to queries about actions that are not mentioned in the agreement.

trans_policy_negative acting on a single *primPolicy* (see process_single_pp_trans_-policy_negative in listing 4.16), first checks that there is a proof for the equality of action_-from_query and action (the action from the agreement). If so, the translation function returns a result of NotPermitted. Otherwise, a result of Unregulated is returned. Note that compared to *trans_policy_positive* we don't check for whether *preRequisite* from the policy holds or not since is only called for policy sets of type *primExclusivePolicySet*.

trans_policy_unregulated acting on a single *primPolicy* (see process_single_pp_-
trans_policy_unregulated in listing 4.15), returns a result of Unregulated.

Listing 4.15: Translation of unregulated policies

```
Definition process_single_pp_trans_policy_unregulated
  (pp: primPolicy)(x:subject)
  (a:asset)(action_from_query: act)  : nonemptylist result :=
 match pp with
  | PrimitivePolicy prq' policyId action ⇒
      (Single (makeResult Unregulated x action_from_query a))
  end.


Fixpoint trans_policy_unregulated
 (e:environment)(x:subject)(p:policy)(a:asset)
 (action_from_query: act){struct p} : nonemptylist result :=

 match p with
    | Policy pp_list ⇒ trans_pp_list_trans_policy_unregulated pp_list x a
    action_from_query
  end.
```

Listing 4.16: Translation of negative policies

```
Definition process_single_pp_trans_policy_negative
  (pp: primPolicy)(x:subject)
  (a:asset)(action_from_query: act)  : nonemptylist result :=
 match pp with
  | PrimitivePolicy prq' policyId action ⇒

      if (eq_nat_dec action_from_query action)
      then
        (Single
          (makeResult NotPermitted x action_from_query a))
      else
        (Single
          (makeResult Unregulated x action_from_query a))

  end.

Fixpoint trans_policy_negative
 (e:environment)(x:subject)(p:policy)(a:asset)
 (action_from_query: act){struct p} : nonemptylist result :=

 match p with
    | Policy pp_list ⇒ trans_pp_list_trans_policy_negative pp_list x a
    action_from_query
  end.
```

Listing 4.17: Translation of positive policies

```
Definition process_single_pp_trans_policy_positive
  (pp: primPolicy)(e:environment)(x:subject)(prin_u:prin)
  (a:asset)(action_from_query: act)  : nonemptylist result :=
 match pp with
   | PrimitivePolicy prq' policyId action ⇒
      if (trans_preRequisite_dec e x prq' (Single policyId) prin_u)
      then (* prin /\ prq /\ prq' *)
       if (eq_nat_dec action_from_query action)
       then
         (Single
           (makeResult Permitted x action_from_query a))
       else
         (Single
           (makeResult Unregulated x action_from_query a))
           else (* prin /\ prq /\ ~prq' *)
        (Single
           (makeResult Unregulated x action_from_query a))
    end.


Definition trans_policy_positive
 (e:environment)(x:subject)(p:policy)(prin_u:prin)(a:asset)
 (action_from_query: act) : nonemptylist result :=
 match p with
     | Policy pp_list ⇒ trans_pp_list_trans_policy_positive pp_list e x prin_u a
    action_from_query
  end.
```

Listing 4.18: Translation of list of policies: positive, negative and unregulated use the same pattern

```
Fixpoint trans_pp_list_trans_policy_<positive | negative | unregulated>
 (pp_list:nonemptylist primPolicy)(e:environment)(x:subject)
 (prin_u:prin)(a:asset)(action_from_query: act){struct pp_list}:=
  match pp_list with
   | Single pp1 ⇒ process_single_pp_trans_policy_<positive | negative | unregulated
    > pp1 e x prin_u a action_from_query
   | NewList pp pp_list' ⇒ app_nonempty
  (process_single_pp_trans_policy_<positive | negative | unregulated> pp e x prin_u
   a action_from_query)
  (trans_pp_list_trans_policy_<positive | negative | unregulated> pp_list' e x
   prin_u a action_from_query)
  end.
```

Listing 4.19: Translation of PIPS and PEPS

```
Definition trans_policy_PIPS
  (e:environment)(prq: preRequisite)(p:policy)(x:subject)
  (prin_u:prin)(a:asset)(action_from_query:act) : nonemptylist result :=

    if (trans_prin_dec x prin_u)
    then (* prin *)
      if (trans_preRequisite_dec e x prq (getId p) prin_u)
      then (* prin /\ prq *)
        (trans_policy_positive e x p prin_u a action_from_query)
      else (* prin /\ ~prq *)
        (trans_policy_unregulated e x p a action_from_query)
    else (* ~prin *)
      (trans_policy_unregulated e x p a action_from_query).

Definition trans_policy_PEPS
  (e:environment)(prq: preRequisite)(p:policy)(x:subject)
  (prin_u:prin)(a:asset)(action_from_query:act) : nonemptylist result :=

    if (trans_prin_dec x prin_u)
    then (* prin *)
      if (trans_preRequisite_dec e x prq (getId p) prin_u)
      then (* prin /\ prq *)
        (trans_policy_positive e x p prin_u a action_from_query)
      else (* prin /\ ~prq *)
        (trans_policy_unregulated e x p a action_from_query)
    else (* ~prin *)
      (trans_policy_negative e x p a action_from_query).
```

Listing 4.20: Translation of Policy Set

```
Fixpoint trans_ps
  (e:environment)(action_from_query:act)(subject_from_query:subject)(asset_from_query
    :asset)
  (ps:policySet)
  (prin_u:prin)(a:asset){struct ps} : nonemptylist result :=

let process_single_ps := (fix process_single_ps (pps: primPolicySet):=

  match pps with
    | PIPS pips =>
      match pips with
        | PrimitiveInclusivePolicySet prq p =>
          (trans_policy_PIPS e prq p subject_from_query prin_u a action_from_query)
      end
    | PEPS peps =>
      match peps with
```

```
      | PrimitiveExclusivePolicySet prq p ⇒
        (trans_policy_PEPS e prq p subject_from_query prin_u a action_from_query)
    end
  end) in

if (eq_nat_dec asset_from_query a)
then (∗ asset_from_query = a ∗)
  match ps with
    | PPS pps ⇒ process_single_ps pps
  end
else (∗ asset_from_query <> a ∗)
    (Single
      (makeResult
        Unregulated subject_from_query action_from_query asset_from_query)).
```

Translation of a *prin* (called *trans_prin* in listing 4.21) takes as input x, the *subject* in question, p, the principal or the *prin*, and proceeds based on whether p is a single subject or a list of subjects. If p is a single subject, s, the *Prop* $x = s$ is returned. Otherwise the disjunction of the translation of the first subject in p (s) and the *rest* of the subjects is returned.

Listing 4.21: Translation of a Prin

```
Fixpoint trans_prin
  (x:subject)(p: prin): Prop :=

  match p with
    | Single s ⇒ (x=s)
    | NewList s rest ⇒ ((x=s) \/ trans_prin x rest)
  end.
```

The translation of a *prerequisite* (called *trans_preRequisite* in listing 4.22) takes as input e, the *environment*, x, the *subject*, *prq*, the *preRequisite* to translate, *IDs*, the set of identifiers (of policies implied by the *prq*), $prin_u$, the agreement's user, and proceeds by case analysis on the structure of the *prerequisite*. A *prerequisite* is either a *TruePrq*, a *Constraint*, a *NotCons*, or a *AndPrqs*.

In listing 4.22 the translation for *TruePrq* is the Prop *True*, the translations for *Constraint*, *NotCons* and *AndPrqs* simply call respective translation functions for corresponding types *constraint* (namely *trans_constraint* and *trans_notCons* and *trans_andPrqs*).

Listing 4.22: Translation of a PreRequisite

```
Definition trans_preRequisite
  (e:environment)(x:subject)(prq:preRequisite)(IDs:nonemptylist policyId)(prin_u:
    prin) : Prop :=

 match prq with
    | TruePrq ⇒ True
    | Constraint const ⇒ trans_constraint e x const IDs prin_u
    | NotCons const ⇒ trans_notCons e x const IDs prin_u
    | AndPrqs prqs ⇒  trans_andPrqs x prq IDs prin_u a
  end.
```

The translation of a *constraint* (called *trans_constraint* in listing 4.23) takes as input *e* the *environment*, *x* the *subject*, *const*, the *constraint* to translate, *IDs*, the set of identifiers (of policies implied by the parent *preRequisite*) and *prin_u*, the agreement's user and proceeds by case analysis on the structure of the *constraint*. A *constraint* is either a *Principal*, a *Count* or a *CountByPrin*. The translation for *Principal* returns the translation function (namely *trans_prin*) for the *prn* (the *prin* that accompanies the *const* constraint). The translation for *Count* and *CountByPrin* return the translation function *trans_count*.

Listing 4.23: Translation of a Constraint

```
Fixpoint trans_constraint
  (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)
  (prin_u:prin){struct const} : Prop :=
 match const with
    | Principal prn ⇒ trans_prin x prn

    | Count n ⇒ trans_count e n IDs prin_u

    | CountByPrin prn n ⇒ trans_count e n IDs prn

  end.
```

The translation of a *NotCons* (called *trans_notCons* in listing 4.24) takes as input *e* the *environment*, *x* the *subject*, *const*, the *constraint* to translate, *IDs*, the set of identifiers (of policies implied by the parent *preRequisite*) and *prin_u*, the agreement's user and proceeds to return the negation of *trans_constraint* (see listing 4.23).

Listing 4.24: Translation of NotCons

```
Definition trans_notCons
  (e:environment)(x:subject)(const:constraint)(IDs:nonemptylist policyId)(prin_u:prin) :
    Prop :=
```

```
~ (trans_constraint e x const IDs prin_u).
```

The translation of a *Count* or a *CountByPrin* (called *trans_count* in listing 4.25) takes as input e the *environment*, n the total number of times the subjects mentioned in $prin_u$ (last parameter) may invoke the policies identified by IDs (third parameter).

To implement the translation for a *Count* or a *CountByPrin* we start by calling an auxiliary function *process_two_lists* that effectively returns a new list composed of pairs of members of the first list and the second list (the cross-product of the two input lists). In the case of *trans_count*, the call is "*process_two_lists IDs prin_u*" which returns a list of pairs of *policyId* and *subject* namely *ids_and_subjects*. *ids_and_subjects* is then passed to a locally defined function *trans_ count_ aux*.

trans_count_aux returns the current count for a single pair of *policyId* and *subject* (the call to *getCount* which looks up the environment e and returns the current count per each *subject* and *policyId*) and for a list of pairs of *policyId* and *subjects*, the addition of *get_count* (for the first pair) and *trans_count_aux*s (for the rest of the pairs) is returned.

A local variable $running_total$ has the value returned by *trans_count_aux*. Finally the proposition $running_total < n$ is returned as the translation for a *Count* or a *CountByPrin*.

Note that the only difference between translations for a *Count* and a *CountByPrin* is the additional *prn* parameter for *CountByPrin* which allows for getting counts for subjects not necessarily the same as $prin_u$, the agreement's user(s).

Listing 4.25: Translation of count

```
Fixpoint trans_count
 (n:nat): Prop :=

 let trans_count_aux
   := (fix trans_count_aux
       (ids_and_subjects : nonemptylist (Twos policyId subject)) : nat :=
    match ids_and_subjects with
      | Single pair1 ⇒ getCount e (right pair1) (left pair1)
      | NewList pair1 rest_pairs ⇒
         (getCount e (right pair1)(left pair1)) +
         (trans_count_aux rest_pairs)
     end) in

 let ids_and_subjects := process_two_lists IDs prin_u in
 let running_total := trans_count_aux ids_and_subjects in
 running_total < n.
```

Chapter 5

Queries

5.1 Introduction

We first mentioned queries in chapter 4 on page 18. Ultimately policy statements describing an agreement will be used to enforce those agreements. To enforce policy agreements, access queries are asked from the policy engine and access is granted or denied based on the answer.

In this chapter we will review our encoding of queries in Coq and Coq representations of other definitions used to prove our decidability results.

5.2 Queries

Queries are tuples of the form $(agreement, s, action, a, e)$. The tuple corresponds to the question of determining whether an agreement implies that a subject s may perform action $action$ on an asset a given the environment e. The Coq representation is listed in listing 5.1.

Listing 5.1: Queries

```
Inductive single_query : Set :=
  | SingletonQuery : agreement → subject → act → asset → environment →
    single_query.
```

5.3 Answering Queries

Answering a query as defined could in general lead to one of four outcomes: error(listing 5.2), permitted(listing 5.3), denied(listing 5.4) and "not applicable" (listing 5.5) defined in [TK06]. Note that e denotes 'an environment being consistent' (see section 4.1) in the following listings.

Listing 5.2: Answerable Queries: Error

$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \implies Permitted(s, act, a) \text{ and}$$
$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \implies \neg Permitted(s, act, a)$$

Listing 5.3: Answerable Queries: Permit

$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \implies Permitted(s, act, a) \text{ and}$$
$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \not\implies \neg Permitted(s, act, a)$$

Listing 5.4: Answerable Queries: Deny

$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \not\implies Permitted(s, act, a) \text{ and}$$
$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \implies \neg Permitted(s, act, a)$$

Listing 5.5: Answerable Queries: Not Applicable

$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \not\implies Permitted(s, act, a) \text{ and}$$
$$\left(\bigwedge[\![agreement]\!]\right) \wedge e \not\implies \neg Permitted(s, act, a)$$

Chapter 6

Formal Proofs and Decidability

6.1   Introduction

We had set out to prove formally that the ACCPL language was decidable. ACCPL started out as a reduced version of Pucella and Weissman's subset of ODRL [PW06] but quickly morphed into a different language, based on the direction of the proof we were developing. An important lesson learned was that the process of designing a certifiable policy language cannot be done in isolation since while certifiability of the language derives the design one way, expressiveness requirements of the language usually drive it in the opposite direction. Certifiability cannot be a post-design activity as the changes it will impose are often fundamental ones.

In the following section we will list the main theorems that we have proved and the supporting theorems that needed to be proved to lead up to the main results. In general however we will not show the actual proofs in the dissertation (with some exceptions) and instead refer the reader to the GitHub repository where the source code for ACCPL, including the proofs, lives [Sis15].

6.2   Decidability of ACCPL

The theorem *trans_agreement_dec* is the declaration of the main decidability result for ACCPL (see listing 6.1). Together with other theorems we will describe below, we have "certified" the decidability of ACCPL. In words, the nonempty list that the agreement translation function *trans_agreement* returns will return a set of results one per each *primPolicy*. Most if not all of the results will be of type Unregulated. In the case where all the results are Unregulated we answer the access query with a result of Unregulated. We show this case indirectly by showing that the set does not contain a Permitted result nor a NotPermitted result which implies in turn the existence of at least one Unregulated result (see 6.2).

In the case where the whole set is not comprised of Unregulated results, we have two mutually exclusive cases. The first case is when the set has a at least one Permitted result;

we answer the access query in this case with a result of Permitted. The second case is when the set has at least one NotPermitted; we answer the access query in this case with a result of NotPermitted (see answer and result types in 4.3).

Listing 6.1: Agreement Translation is Decidable

```
Theorem trans_agreement_dec2:
 ∀
 (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),


 (isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query))
\/

 (isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query))
\/

 (~(isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query)) /\
  ~(isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query))).
```

Listing 6.2: No Permitted and No NotPermitted Implies Unregulated

```
Theorem trans_agreement_not_NotPerm_and_not_Perm_implies_Unregulated_dec:
 ∀
  (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),

  (~(isResultInQueryResult
    (Result Permitted subject_from_query action_from_query asset_from_query)
     (trans_agreement e ag action_from_query subject_from_query asset_from_query))
    /\

   ~(isResultInQueryResult
    (Result NotPermitted subject_from_query action_from_query asset_from_query)
     (trans_agreement e ag action_from_query subject_from_query asset_from_query)))
     →

  (isResultInQueryResult
```

33

```
  (Result Unregulated subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query)).
```

The next three theorems establish the mutual exclusivity of Permitted and NotPermitted results. The *trans_agreement_perm_implies_not_notPerm_dec* (see listing 6.3) says that the existence of a Permitted result in the set of results returned by *trans_agreement*, excludes a NotPermitted result while *trans_agreement_NotPerm_implies_not_Perm_dec* (see listing 6.4) shows the inverse. Finally the proof for *trans_agreement_not_Perm_and_NotPerm_at_once* establishes that both Permitted and NotPermitted results cannot exist in the same set returned by *trans_agreement* (see listing 6.5). This result also establishes the fact that in ACCPL conflicts are not possible.

Listing 6.3: Permitted result Implies no NotPermitted result

```
Theorem trans_agreement_perm_implies_not_notPerm_dec:
 ∀
  (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),

  (isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
    (trans_agreement e ag action_from_query subject_from_query asset_from_query))
   →

 ~(isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
    (trans_agreement e ag action_from_query subject_from_query asset_from_query)).
```

Listing 6.4: NotPermitted result Implies no Permitted result

```
Theorem trans_agreement_NotPerm_implies_not_Perm_dec:
 ∀
  (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),

  (isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
    (trans_agreement e ag action_from_query subject_from_query asset_from_query))
   →

 ~(isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
    (trans_agreement e ag action_from_query subject_from_query asset_from_query)).
```

Listing 6.5: Permitted and NotPermitted results mutually exclusive

```
Theorem trans_agreement_not_Perm_and_NotPerm_at_once:
 ∀
 (e:environment)(ag:agreement)(action_from_query:act)
```

```
  (subject_from_query:subject)(asset_from_query:asset),


 ~((isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query))
/\
 (isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query))).
```

The proof for the next theorem *trans_agreement_not_NotPerm_and_not_Perm_implies_Unregulated_dec* shows that in the case where neither a Permitted nor NotPermitted result exists in the set returned by *trans_agreement*, there does exist at least one Unregulated result (see listing 6.6).

Listing 6.6: Not (Permitted and NotPermitted) Implies Unregulated

```
Theorem trans_agreement_not_NotPerm_and_not_Perm_implies_Unregulated_dec:
 ∀
   (e:environment)(ag:agreement)(action_from_query:act)
   (subject_from_query:subject)(asset_from_query:asset),

   (~(isResultInQueryResult
     (Result Permitted subject_from_query action_from_query asset_from_query)
       (trans_agreement e ag action_from_query subject_from_query asset_from_query))
     /\
   ~(isResultInQueryResult
     (Result NotPermitted subject_from_query action_from_query asset_from_query)
       (trans_agreement e ag action_from_query subject_from_query asset_from_query)))
     →
 (isResultInQueryResult
     (Result Unregulated subject_from_query action_from_query asset_from_query)
      (trans_agreement e ag action_from_query subject_from_query asset_from_query)).
```

6.3   Proofs Hierarchy

All the theorems shown in the listings of section 6.2, specify *trans_agreement* as the top translation function that returns a set of results one per the *primPolicy* in the agreement. We have developed the proofs in a modular fashion such that the hierarchy of the proofs follow the structure of ACCPL language constructs very closely. For example, we have a version of the main theorem 6.1 for the translation functions for a *policySet*, a

primInclusivePolicySet, a *primExclusivePolicySet*, etc. See listings 6.7 for the version for the translation function for a *policySet*.

Listing 6.7: PolicySet Translation is Decidable

```
Theorem trans_ps_dec:
 ∀
 (e:environment)(action_from_query:act)(subject_from_query:subject)(asset_from_query
    :asset)
 (ps:policySet)
 (prin_u:prin)(a:asset),

 (isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
   (trans_ps e action_from_query subject_from_query asset_from_query ps prin_u a))
\/

 (isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
   (trans_ps e action_from_query subject_from_query asset_from_query ps prin_u a))
\/

 (isResultInQueryResult
   (Result Unregulated subject_from_query action_from_query asset_from_query)
   (trans_ps e action_from_query subject_from_query asset_from_query ps prin_u a)).
```

While each proof builds on top of a previously proven decidability theorem based on the structure of the language construct, the pattern of the propositions to prove doesn't always stay the same. For example, the proof for *trans_ps_dec* uses the proofs for the following two theorems: *trans_policy_PIPS_dec* and *trans_policy_PEPS_dec*. However the statement of decidability is different for *trans_policy_PIPS_dec* as the set returned by this translation function will not contain a NotPermitted result; either a Permitted or an Unregulated result will be returned in this case (see listing 6.8).

Listing 6.8: Primitive Inclusive PolicySet Translation is Decidable

```
Theorem trans_policy_PIPS_dec:
 ∀
 (e:environment)(prq: preRequisite)(p:policy)(subject_from_query:subject)
 (prin_u:prin)(a:asset)(action_from_query:act),

 (isResultInQueryResult
   (Result Permitted subject_from_query action_from_query a)
   (trans_policy_PIPS e prq p subject_from_query prin_u a action_from_query))
\/

 (isResultInQueryResult
   (Result Unregulated subject_from_query action_from_query a)
```

```
(trans_policy_PIPS e prq p subject_from_query prin_u a action_from_query)).
```

6.3.1   Decidable As An Inductive Predicate

We encode the cases discussed in section 6.2 in an inductively defined predicate called *decidable* in listing 6.9. Recall that the nonempty list that the agreement translation function *trans_agreement* returns will return a set of results one per each *primPolicy*.There are three cases. First when there is a NotPermitted result in the set, second when there is a Permitted result in the set, and third when there are no Permitted and no NotPermitted results in the set. We encode these three cases in three constructors for the type decidable. The constructors are: *Denied*, *Granted* and *NotApplicable*.

Listing 6.9: Inductive Predicate decidable

```
Inductive decidable : environment → agreement → act → subject → asset → Prop :=

    | Denied : ∀
 (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),

 (isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
     (trans_agreement e ag action_from_query subject_from_query asset_from_query))
 → decidable e ag action_from_query subject_from_query asset_from_query

    | Granted : ∀
 (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),
 (isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
     (trans_agreement e ag action_from_query subject_from_query asset_from_query))
 → decidable e ag action_from_query subject_from_query asset_from_query

    | NonApplicable : ∀
 (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),

 ~(isResultInQueryResult
   (Result Permitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query))
    →
 ~(isResultInQueryResult
   (Result NotPermitted subject_from_query action_from_query asset_from_query)
   (trans_agreement e ag action_from_query subject_from_query asset_from_query))
 → decidable e ag action_from_query subject_from_query asset_from_query.
```

We have declared (and proven) the decidability of ACCPL using the "decidable" type as seen in listing 6.10.

Listing 6.10: Decidability Of Agreement Translation

```
Theorem decidable_trans_agreement:
 ∀
 (e:environment)(ag:agreement)(action_from_query:act)
 (subject_from_query:subject)(asset_from_query:asset),
   decidable e ag action_from_query subject_from_query asset_from_query.
```

6.4 ACCPL in the Landscape of Policy-based Policy Languages

Tschantz and Krishnamurthi [TK06] argue for the need for formal means to compare and contrast policy-based access-control languages and they present a set of properties to analyze the behaviour of policies in light of additional and/or explicit environmental facts, policy growth and policy decomposition. Tschantz and Krishnamurthi apply their properties to two core policy languages and compare the results. One of the core policy languages they present is a simplified XACML and the other is Lithium [HW08]. In the following sections we will describe Tschantz and Krishnamurthi's properties and evaluate ACCPL with respect to these properties.

6.5 Reasonability Properties

The properties Tschantz and Krishnamurthi [TK06] present revolve around three questions:

1. How decisions change when the environment has more facts

2. Impact of policy growth on how decisions may change

3. How amenable policies are to compositional reasoning

Tschantz and Krishnamurthi [TK06] use a simple policy as a motivating example (which they attribute to [HW08]) that we will adapt and reuse here. The example will help to review the range of interpretations possible and ultimately the range of policy language design choices and where ACCPL is on that range. The policy in English is as follows:

6.5.1 Basic Policy and Query

- If the subject has a subscription to the Gold bundle, then permit that subject to print TheReport.

- If the subject has a subscription to the Basic bundle, then do not permit that subject to print TheReport.

- If the subject has no subscription to the Gold bundle, then permit that subject to play Terminator.

Consider a query and a fact in English as: A subscriber to the Basic bundle requests to play the movie Terminator. Let's extract the fact that the subject is a subscriber to the Basic bundle into a separate term and make it part of the environment. The basic policy, the query in question and the environment are all expressed as a (pseudo) logic formula and listed in the listing 6.11. Should the access query be granted?

Listing 6.11: Basic Policy

```
p1 = GoldBundle(s) → Permitted(s, print, TheReport)
p2 = BasicBundle(s) → NotPermitted(s, print, TheReport)
p3 = ~GoldBundle(s) → Permitted(s, play, Terminator)
q1 = (s, play, Terminator)
e1 = BasicBundle(s)
```

According to Tschantz and Krishnamurthi [TK06] at least three different interpretations are possible. We will list the possibilities and then evaluate how ACCPL would answer the query.

6.5.2 Possible Interpretations

- Implicit: Grant access because of $p3$. This interpretation assumes that if a subject has a subscription to the Basic bundle, the subject does not have a subscription to the Gold bundle so the assumption is that there is a proof for $GoldBundle(s)$. This is an example of "implicit" knowledge and "closed world assumptions".

- Explicit: The policy does not apply to the query. In other words, the query is "NotApplicable". $p1$ and $p2$ don't apply since the asset and the action in the request don't match the ones in $p1$ and $p2$. $p3$ does not apply either since no proof exists to show that $GoldBundle(s)$. This interpretation is similar to how ACCPL would answer the query.

- Implicit with automatic proof capability: grants access to the query by automatically proving that $GoldBundle(s)$. Under this interpretation, proof by contradiction is used to establish $GoldBundle(s)$ since if $GoldBundle(s)$, then $p1$ and $p2$ would contradict.

6.6   Policy Based Access-Control Languages

Tschantz and Krishnamurthi [TK06] define an access-control policy language as a tuple $L = (P, Q, G, N, \langle\langle . \rangle\rangle)$ where P is a set of policies, Q is a set of requests or queries, G is the granting decisions, N is the non-granting decisions, $\langle\langle . \rangle\rangle$ is a function taking a policy $p \in P$ to a relation between Q and $G \cup N$ where $G \cap N = \emptyset$. Let D represent $G \cup N$. Policy p assigns a decision of $d \in D$ to the query $q \in Q$. L also defines a partial order on decisions such that $d <= d'$ if either $d, d' \in N$ or $d, d' \in G$ or $d \in N$ and $d' \in G$, so for example $Denied <= Granted$ in ACCPL.

6.7   Types of Decisions

Policy based access-control languages typically support at least the Permit and Deny decisions. When a decision for a query is not granted, one design choice for a language is to return a Deny. However in this case Deny stands for "not Permit". It is possible to have cases when the policy truly doesn't specify either Permit or Deny. In such cases arbitrarily returning the decision of Deny makes it difficult to compose policies and in fact, an explicit decision of NonApplicable is warranted. Some languages may decide to only have Permit decisions. In such languages lack of a Permit decision for a query signifies a Deny so Deny decisions are not explicit. Although the policies of these languages may be more readable than those with more explicit decisions but they result in ambiguity on whether a Deny decision was really intended or not. Finally some languages define an explicit decision of "Error" for cases such as when both Permit and Deny decisions are reached for the same query. An explicit Error decision is preferable to undefined behaviour because it can lead to improvements to policies and/or how the queries are built.

6.8   Policy Combinators

Primitive policies are often developed by independent entities within an organization therefore it is natural to want a way of combining them into a policy for the whole organization using policy combinators. Some languages such as ODRL have nested sub-policies inside other policies. In such cases one design decision could be to allow for different combinators for different layers to make the language more expressive. Some examples of policy combinators are: conjunction, disjunction, exclusive disjunction, Permit override (if any of the primitive policies returns a Permit, return only Permit), Deny override (if any of the primitive policies returns a Deny, return only Deny).

6.9   ACCPL

When designing ACCPL and to make the decidability proofs possible, we sought out all the different combinations and cases possible (through interaction with the proof process)

and explicitly enumerated them in the translation functions. Once we made explicit all the different cases where a decision was possible, we assigned Permitted and NotPermitted decisions (equivalent to Permit and Deny decisions in this chapter) to two specific cases while every other enumerated case was rendered with an Unregulated (not applicable) decision. Explicit enumeration of all the different cases has been certified to be complete by the fact that we have the decidability results. For future work, one can modify the decision assignments to different cases, deciding for the semantics to give more Permit decisions for example or more Deny decisions or other changes in the same vein. One can also remove the Unregulated decisions from the translation functions and assign instead Permit or Deny depending on whether we want to make the language more permissive or more prohibitive. In this case however the current decidability theorem will no longer hold and the Unregulated case would have to be removed from the statement of the theorem and the proof to be modified to look like 6.12.

Listing 6.12: Agreement Translation is Decidable (for Permitted or NotPermitted only)

```
Theorem trans_agreement_dec:
 ∀
 (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset),

 (isResultInQueryResult
    (Result Permitted subject_from_query action_from_query asset_from_query)
    (trans_agreement e ag action_from_query subject_from_query asset_from_query))
\/
 (isResultInQueryResult
    (Result NotPermitted subject_from_query action_from_query asset_from_query)
    (trans_agreement e ag action_from_query subject_from_query asset_from_query)).
```

6.10   Reasonability Properties for ACCPL

6.10.1   Deterministic and Total

According to Tschantz and Krishnamurthi [TK06] a language L is defined to be deterministic if $forall\ p \in P,\ forall\ q \in Q,\ forall\ d, d' \in D,\ q \langle\langle p \rangle\rangle\, d \land q \langle\langle p \rangle\rangle\, d' \implies d = d'$. A language L is total if $forall\ p \in P,\ forall\ q \in Q,\ exists\ d \in D$, s.t. $q \langle\langle p \rangle\rangle\, d$.

We declare (and prove) the determinism of ACCPL in the theorem listed in 6.13. ACCPL is proved total by the fact that it is decidable (see 6.10).

Listing 6.13: Agreement Translation is Deterministic

```
Theorem trans_agreement_deterministic:
 ∀
 (e:environment)(ag:agreement)(action_from_query:act)
  (subject_from_query:subject)(asset_from_query:asset)(d: answer),
```

```
(isResultInQueryResult
 (Result d subject_from_query action_from_query asset_from_query)
 (trans_agreement e ag action_from_query subject_from_query asset_from_query)) →

(isResultInQueryResult
 (Result d subject_from_query action_from_query asset_from_query)
 (trans_agreement e ag action_from_query subject_from_query asset_from_query)).
```

6.10.2   Safety

As discussed in section 6.5.2 policy interpretations could be implicit or explicit. Explicit interpretations distinguish between unknown information and information known to be absent. The explicit approach promotes verbose policies and requests and usually leads to too many Unregulated or NotApplicable decisions (in the absence of the right environmental facts). However the verbosity could be used to direct what additional facts are needed to make a granting and denying decision. Under the second interpretation in 6.5.2 while the decision is NotApplicable, the system can direct the entity making the request provide a proof of $GoldBundle(s)$ in order for a granting decision to be rendered.

Implicit interpretations may result in granting unintended access to protected resources because no explicit facts need be mentioned in the environment. Recall that the implicit approach works by making assumptions when facts are not explicitly mentioned. Such unintended assumptions lead to unintended permissions being granted.

Tschantz and Krishnamurthi [TK06] define safety as $forall\ p \in P,\ forall\ q, q' \in Q,$ $q \leq q' \implies q \langle\langle p \rangle\rangle d \leq q \langle\langle p \rangle\rangle d'$ which intuitively says that policies written in a safe language L, will not result in the leakage of permissions to unintended subjects. For example, incomplete information may unintentionally cause an access permission to be granted to an intruder.

Could we conclude that an access-control policy language with the explicit approach such as ACCPL is also safe? in the design of ACCPL and the translation functions in particular we only grant or deny access in specific cases. We grant access if the asset and the action in query are found in the agreement and $trans\_prin$ has a proof and $trans\_preRequisite$ for both policy set and policy have proofs. We deny an access request if the policy set is an exclusive policy set ($primExclusivePolicySet$) and $trans\_prin$ does not hold. In all other cases we render the decision of Unregulated. In other words, in the absence of necessary facts (proofs), we give the least decision (e.g. Unregulated), when there are more facts available in the environment (such as $trans\_preRequisite$ holding), we grant either deny or permit decisions. However when all the facts hold, corresponding to the greatest query in the definition of safety, we render a granting decision which is also the greatest decision in the lattice of decisions: Unregulated ≤ NotPermitted ≤ Permitted.

6.10.3 Independent Composition

As discussed in section 6.5.2, the third interpretation works by taking into account not only the third policy, $p3$, but also $p1$ and $p2$ to render a decision of granted. However each of the policies in isolation would result in the NotApplicable decision.

Taking into account all policies and rendering a decision that is not different from combining the decisions reached by each primitive policy in isolation, is a property called "Independent Composition" by Tschantz and Krishnamurthi [TK06]. The third interpretation in section 6.5.2 is clearly not have this property. Formally, a language L has the independent composition property, if $forall\ p \in P,\ forall\ q \in Q,\ forall\ d, d^* \in D$,
$\boxplus (q \langle\langle p_1 \rangle\rangle d_1,\ q \langle\langle p_2 \rangle\rangle d_2,\ ...,\ q \langle\langle p_n \rangle\rangle d_n) = q \langle\langle (\oplus p_1, p_2, ..., p_n) \rangle\rangle d^*$, where $\oplus$ is a composition operator defined in the language and $\boxplus$ is the decision composition operator. For this definition to be well-defined, there is a requirement that complete environmental facts be available for each individual policy so that a decision can be rendered for each primitive policy with respect to the request, in isolation.

Does ACCPL have the independent composition property? As covered in chapter 4 and subsection 4.4.1 the translation functions for ACCPL, starting with $trans\_agreement$ all return sets of $result$s. The set of $result$s returned by the agreement translation function will have a result per $primPolicy$ (type of primitive policies) in the agreement. Recall that results containing a Permitted answer and a NotPermitted answer are mutually exclusive and the existence of a single (or more) Permitted result in the set makes the whole set "Permitted" whereas the existence of a single (or more) NotPermitted result makes the whole set "NotPermitted". When no Permitted or NotPermitted is seen in the set, we only have Unregulated results which would make the whole set "Unregulated". This is the decision combination strategy we have designed for ACCPL; the policy combination strategy ACCPL simply preserves all individual decisions and returns the whole set for rendering a global decision. Below we consider the three different possible decisions that the decision combination strategy may reach and review whether the result of any individual policy may change as a result.

1. Global decision is Permitted. Only possible results are Permitted or Unregulated. In this case, only the Unregulated results have been modified by the decision combination strategy. Permitted results coming from any individual primitive policy will have been preserved.

2. Global decision is NotPermitted. Only possible results are NotPermitted or Unregulated. In this case, similar to the first case, only the Unregulated results have been modified by the decision combination strategy. NotPermitted results coming from any individual primitive policy will have been preserved.

3. Global decision is Unregulated. Only possible results are Unregulated. In this case, results coming from any individual primitive policy will have been preserved.

We argue that ACCPL has the independent composition property with respect to granting and denying policies only as detailed in the previous paragraphs.

6.10.4 Monotonicity of Policy Combinators

Tschantz and Krishnamurthi [TK06] discuss the notion of monotonicity of policy combinators. Intuitively, in a language with a policy combinator that is monotonic, adding another primitive policy, cannot change the combined decision to go from grating to non-granting. In ACCPL, NotApplicable results are only possible if $trans\_prin$ does not hold. Assume the combined (global) decision is already Permitted for a certain agreement. In the case when $trans\_prin$ does not hold, adding a primitive policy to the agreement will not change the Permitted result to a NotPermitted result, since a primitive policy is composed of a $preRequisite$ and an act only. In the case where $trans\_prin$ does hold, there are two cases. Either the $preRequisite$ of the added policy holds and the decision is another Permitted result, or the $preRequisite$ of the added policy does not hold, and the decision is an Unregulated which won't affect the combined decision.

Chapter 7

Examples

7.1 Introduction

In this chapter we will take a tour of the syntax and semantics we have so far developed by examining some example agreements. In the following we will start by reviewing some of the examples used in [PW06].

Ultimately the goal of specifying all the syntax and semantics is to declare some interesting theorems about policy expressions and proving them in Coq. However we will first start with some specific propositions/theorems about these examples to get a feel for how proofs are done in Coq.

7.2 Agreement 2.1

Consider example 2.1 (from [PW06]) where the *policySet* is a *AndPolicySet* with $p1$ and $p2$ as the individual *policySet*s. Let $p1$ be defined as $Count[5] \rightarrow print$ and $p2$ as $and[Alice, Count[2]] \rightarrow print$.

The agreement is that the asset *The Report* may be printed a total of five times by either *Alice* or *Bob*, and twice more by Alice. So if Alice and Bob have used policy $p1$ to justify their printing of the asset m_{p1} and n_{p1} times, respectively, then either may do so again if $m_{p1} + n_{p1} < 5$. If they have used $p2$ to justify their printing of the asset m_{p2} and n_{p2} times, respectively, then only Alice may do so again if $m_{p2} + n_{p2} < 2$. Note that since Bob doesn't meet the prerequisite of being Alice, n_{p2} is effectively 0, so we have $m_{p2} < 2$ as the condition for Alice being able to print again (Alice does meet the prerequisite of being Alice).

Listing 7.1: Agreement 2.1 (as used in [PW06])

```
agreement
 for {Alice, Bob}
 about TheReport
 with and [p1, p2].
```

The Coq version of the agreement 2.1 (listing 7.1) and its sub-parts is listed below. It is best to start with the agreement itself called $A2.1$ in the listing and compare to the agreement 2.1 listed in 7.1.

Listing 7.2: Agreement 2.1 in Coq

```
Definition p1A1:policySet :=
  PrimitivePolicySet
    TruePrq
    (PrimitivePolicy (Constraint (Count 5)) id1 Print).

Definition p2A1prq1:preRequisite := (Constraint (Principal (Single Alice))).
Definition p2A1prq2:preRequisite := (Constraint (Count 2)).

Definition p2A1:policySet :=
  PrimitivePolicySet
    TruePrq
    (PrimitivePolicy (AndPrqs (NewList p2A1prq1 (Single p2A1prq2))) id2 Print).

Definition A2.1 := Agreement (NewList Alice (Single Bob)) TheReport
                (AndPolicySet (NewList p1A1 (Single p2A1))).
```

7.3 Agreement 2.5

Consider example 2.5 (from [PW06]) where the *policySet* is a *PrimitivePolicySet* with a *Count* constraint as prerequisite and a *AndPolicy* as the policy. The *AndPolicy* is the conjunction of two *PrimitivePolicy*s. Both policies have prerequisites of type *ForEachMember* with actions *display* and *print* respectively. The *prin* component for both *ForEachMember*s is *Alice, Bob*, whereas the constraint for the first *ForEachMember* is *Count*[5] and for the second is *Count*[2].

Listing 7.3: Agreement 2.5 (as used in [PW06])

```
agreement
 for {Alice, Bob}
 about ebook
 with Count [10] → and [forEachMember[{Alice, Bob}; Count[5]] ⇒id1 display,
                       forEachMember[{Alice, Bob}; Count[1]] ⇒id2 print].
```

The Coq version of the agreement 2.5 (listing 7.3) and its sub-parts is listed below. See agreement 2.5 listed in 7.3 for comparison.

The agreement is that the asset *ebook* may be displayed up to five times by Alice and Bob each, and printed once by each. However the total number of actions (either *display* or *print*) justified by the two policies by either Alice and Bob is at most 10.

Listing 7.4: Example 2.5

```
Definition tenCount:preRequisite := (Constraint (Count 10)).
Definition fiveCount:constraint := (Count 5).
Definition oneCount:constraint := (Count 1).

Definition prins2_5 := (NewList Alice (Single Bob)).
Definition forEach_display:preRequisite := ForEachMember prins2_5 (Single fiveCount)
    .
Definition forEach_print:preRequisite := ForEachMember prins2_5 (Single oneCount).

Definition primPolicy1:policy := PrimitivePolicy forEach_display id1 Display.
Definition primPolicy2:policy := PrimitivePolicy forEach_print id2 Print.

Definition policySet2_5:policySet :=
  PrimitivePolicySet tenCount (AndPolicy (NewList primPolicy1 (Single primPolicy2))).

Definition A2.5 := Agreement prins2_5 ebook policySet2_5.
```

Chapter 8

Some Simple Theorems

8.1 Introduction

In this chapter we will declare and prove some very simple theorems about the examples from chapter 7. This simple introduction is only meant to give us a feel for how theorems are stated in Coq and how proofs are constructed using Coq *tactics*.

As mentioned earlier, propositions are types in Coq whose type is the sort *Prop*. Any term t whose type is a proposition is a proof term or, for short, a proof. A *Hypothesis* is a local declaration $h : P$ where h is an identifier and P is a proposition. An *Axiom* is similar to a hypothesis except it is declared at the global scope and so it is always available.

A *Theorem* or *Lemma* is stated by giving an identifier whose type is a proposition ([BC04]). The proposition is the statement of the theorem or lemma. It must be followed by a proof. Keywords "Hypothesis", "Axiom" and "Theorem" or "Lemma" are used in each case respectively.

To build a proof in Coq the user states the proposition to prove; this is called a goal to be proved or discharge, along with some hypothesis that makes up the local context. The user then uses commands called tactics to manipulate the local context and to decompose the goal into simpler goals. The goal simplification into sub-goals will continue until all the sub-goals are solved.

In listing 8.1 we have declared a theorem called *example*1 and the corresponding proposition $for all x : nat, x < x + 1$.

Note that the notation $P : T$ is also used to declare program P has type T. This duality of notation is due to Curry-Howard isomorphism which relates the two worlds of type theory and structural logic together [BC04]. Once the Theorem has been declared Coq displays the proposition to be proved under a horizontal line written ——, and displays the context of local facts and hypothesis, if any, above the horizontal line. At this point one can enter proof mode by using *Proof.* upon which Coq is ready to accept tactics. Entering tactics that can break the stated goal (under the horizontal line) into one or more sub-goals is how one progresses until no goals left at which point Coq responds with "No more subgoals" [Ber10].

Listing 8.1: Proof Example

```
Theorem example1: ∀ x:nat, x < x + 1.
```

In the following listings, for the sake of completeness of the presentation, we will include the Coq commands that complete the proof of the respective theorem. We will not however explain the individual commands further here. Showing the commands in the listings, is meant as an indication of the size of the proof in terms of lines of Coq script.

8.2   Theorem One

In listing 8.3 we define a *policySet* with a *constraint* such that if *Alice* has used the policy with *id*1 to justify her printing a_1 times, she may do so again if $a_1 < 5$. The agreement *AgreeCan* simply links the asset *TheReport* with the subject *Alice* and the *policySet* previously defined.

We capture the fact that *Alice* has used the policy with *id*1 to justify her printing 2 times in an environment called *eA*1. Recall that environments are defined to be non-empty lists of $count_e quality$ objects (see listing 4.1).

We also declare a hypothesis H with the proposition that results from the translation of the agreement (see definition of $trans_a greement$ in listing 4.14) and the environment. The proposition can be shown in Coq after some clean-up (e.g. replaced 101 by Alice) and using the form *Eval compute* :

Listing 8.2: Hypothesis for Theorem One

```
∀ x : subject, x = Alice /\ True → 2 < 5 → Permitted x Print TheReport.
```

The theorem *One* that we are going to prove is trivial but nonetheless in English it states that Alice is Permitted to Print TheReport. The proof comes after the command 'Proof.' and ends with 'Qed'.

Listing 8.3: Theorem One

```
Definition psA1:policySet :=
 PrimitivePolicySet
   TruePrq
   (PrimitivePolicy (Constraint (Count 5)) id1 Print).

Definition AgreeCan := Agreement (Single Alice) TheReport psA1.

Definition eA1 : environment :=
 (SingleEnv (make_count_equality Alice id1 2)).

Hypothesis H: trans_agreement eA1 AgreeCan.

Theorem One: Permitted Alice Print TheReport.
```

```
Proof. simpl in H. apply H. split. reflexivity. auto.  omega. Qed.
```

8.3   Theorem Two

In listing 8.5 we define an exclusive policy set *policySet* containing a policy *pol* that allows printing. The agreement *AgreeA5* includes the exclusive policy set to express that Bob may print *LoveAndPeace*. However any subject that is not the agreement's user (e.g. Bob) is forbidden from printing *LoveAndPeace*.

Notice that due to the fact that environments are defined as non-empty lists, we have added a Null count to it (see *eA5*). We continue to capture the relevant facts from the environment and the agreement through defining a hypothesis (e.g. *H*). The hypothesis is shown in listing 8.5.

Listing 8.4: Hypothesis for Theorem Two

```
∀ x : subject,
      (x = Bob /\ True → True → Permitted x Print LoveAndPeace) /\
      ((x = Bob → False) → Permitted x Print LoveAndPeace → False).
```

Theorem $T1\_A5$ states the exclusivity of the policy set, namely that any subject that is not Bob is not permitted to print the asset LoveAndPeace. Theorem $T2\_A5$ uses $T1\_A5$ to prove Alice is not permitted to print the asset.

Listing 8.5: Theorem Two

```
Definition prin_bob := (Single Bob).
Definition pol:policy := PrimitivePolicy TruePrq id3 Print.
Definition pol_set:policySet := PrimitiveExclusivePolicySet TruePrq pol.
Definition AgreeA5 := Agreement prin_bob LoveAndPeace pol_set.
Definition eA5 : environment := (SingleEnv (make_count_equality NullSubject NullId 0))
    .


Hypothesis H: trans_agreement eA5 AgreeA5.


Theorem T1_A5: ∀ x, x<>Bob → ~Permitted x Print LoveAndPeace.
Proof. simpl in H. apply H. Qed.

Theorem T2_A5: ~Permitted Alice Print LoveAndPeace.
Proof. simpl in H. apply T1_A5. apply not_eq_S. omega. Qed.


End A5.
```

8.4   Theorem Three

In listing 4.1 we defined environments as non-empty lists of $count_e quality$ objects which are in turn defined as counts per each subject, policy-id pair. These count formulas represent how many times each policy has been used to justify an action by a subject wrt a policy (specified by the policy id) and semantically it makes sense that they are unique in time. When and if two $count_e quality$ objects with the same subject and policy id refer to different counts, we say we have *inconsistent* count formulas. The listing 8.6 defines the binary predicate *inconsistent*.

Listing 8.6: Inconsistent Count Formulas

```
Definition inconsistent (f1 f2 : count_equality) : Prop :=
  match f1 with (CountEquality s1 id1 n1) ⇒
    match f2 with (CountEquality s2 id2 n2) ⇒
      s1 = s2 → id1 = id2 → n1 <> n2
    end
  end.
```

Next we would like to expand the notion of inconsistency to more than two count formulas. We first define a predicate over a count formula and an environment as in listing 8.7. If the environment is a singleton then we just compare the two count formulas for inconsistency, else we build the disjunction of the inconsistency between the count formula on one hand and the head of the environment and the rest of the environment, respectively.

Listing 8.7: Inconsistent Count Formula And Environment

```
Fixpoint formula_inconsistent_with_env (f: count_equality)
                 (e : environment) : Prop :=
  match e with
    | SingleEnv g ⇒  inconsistent f g
    | ConsEnv g rest ⇒ (inconsistent f g) \/ (formula_inconsistent_with_env f rest)
  end.
```

Finally we define a new inductive data type that represents *consistent* environments (see listing 8.8). An environment is consistent if it is a singleton count formula, if it consists of only two consistent count formulas and finally if the environment consists of a consistent environment and the consistent composition of a count formula and the consistent environment (see constructor *consis_more*.

Listing 8.8: Inconsistent Environment

```
Inductive env_consistent : environment → Prop :=
| consis_1 : ∀ f, env_consistent (SingleEnv f)
| consis_2 : ∀ f g, ~(inconsistent f g) → env_consistent (ConsEnv f (SingleEnv g))
| consis_more : ∀ f e,
```

```
  env_consistent e → ~(formula_inconsistent_with_env f e) → env_consistent (
    ConsEnv f e).
```

We will now pose several small theorems about consistency of count formulas and environments and provide proofs for them (see listing 8.9).

Listing 8.9: Inconsistent Environment

```
Theorem f1_and_f2_are_inconsistent: inconsistent f1 f2.
Proof.
unfold inconsistent. simpl. omega. Qed.


Theorem f1_and___env_of_f2_inconsistent: formula_inconsistent_with_env f1 (
    SingleEnv f2).
Proof.
unfold formula_inconsistent_with_env. apply f1_and_f2_are_inconsistent. Qed.


Theorem two_inconsistent_formulas_imply_env_inconsistent:
 ∀ f g, inconsistent f g → ~env_consistent (ConsEnv f (SingleEnv g)).
Proof.
intros. unfold not. intros H'.
inversion H'. intuition. intuition. Qed.


Theorem e2_is_inconsistent: ~env_consistent e2.
Proof.
apply two_inconsistent_formulas_imply_env_inconsistent.
apply f1_and_f2_are_inconsistent. Qed.


Theorem env_consistent_implies_two_consistent_formulas:
 ∀ (f g: count_equality),
   env_consistent (ConsEnv f (SingleEnv g))→ ~inconsistent f g.
Proof.
intros. inversion H. exact H1. intuition. Qed.


Theorem two_consistent_formulas_imply_env_consistent:
 ∀ (f g: count_equality),
   ~inconsistent f g → env_consistent (ConsEnv f (SingleEnv g)).
Proof.
intros. apply consis_2. exact H. Qed.

Theorem env_inconsistent_implies_two_inconsistent_formulas:
 ∀ (f g: count_equality),
   ~env_consistent (ConsEnv f (SingleEnv g))→ inconsistent f g.
```

```
Proof.
induction f.
induction g.
unfold inconsistent.
intros.
subst.
generalize (dec_eq_nat n n0).
intro h; elim h.
intro; subst.
elim H.
apply consis_2.
unfold inconsistent.
intro.
assert (s0=s0); auto.
assert (p0=p0); auto.
specialize H0 with (1:=H1) (2:=H2).
elim H0; auto.
auto.
Qed.


Theorem same_subjects_policyids_different_counts_means_inconsistent : ∀ (s1 s2:
    subject),
         ∀ (id1 id2: policyId),
         ∀ (n1 n2: nat),

 (s1 = s2 /\ id1 = id2 /\ n1 <> n2) →
 inconsistent (CountEquality s1 id1 n1) (CountEquality s2 id2 n2).

Proof.
intros. unfold inconsistent. intros. intuition. Qed.
```

Chapter 9

Proposed Future Work

9.1 Summary

We started off by looking at DRELs and specifically at ODRL and considered its formal semantics as captured by [PW06]. We presented the encodings and semantics of the constructs for a significant subset of ODRL in Coq and then defined what queries looked like and what the decision problem was in this context. We have also encoded the decision algorithms as presented in [PW06] in Coq in order to perform formal verification of theorems of interest. We noted the common thread between RELs and policy languages for access control systems such as those between ODRL and SELinux. We discussed the goal of generalizing the concept of a policy language from strictly representing subsets of ODRL, to representing (subsets of) both ODRL and SELinux policy languages with the goal of applying the decision algorithms to both types of policies, in a unified manner.

9.2 Machine-Checked Proof of Decidability of Queries

By defining formal semantics for ODRL authors of [PW06] were able to show some important results. First result is that answering the question of whether a set of ODRL statements imply a permission, denial or other possibilities is decidable and also that its complexity is NP-hard.

The authors of [PW06] then prove that by removing the construct *not[policySet]* from ODRL's syntax answering the same query remains decidable and efficient (polynomial time complexity).

We will prove equivalent results as above starting with the decidability result of answering a query in ODRL0 (which does not include *not[policySet]*). The theorem in listing 9.1 states that for all environments, all single agreements, all subjects, all actions and all assets, either permission is granted, permission is denied, permission is unregulated or query is inconsistent.

Listing 9.1: Environments and Counts

```
Theorem queriesAreDecidable: ∀ (e:environment),
              ∀ (agr: agreement),
              ∀ (s:subject),
              ∀ (action:act),
              ∀ (a:asset),

(permissionGranted e [agr] s action a) \/
(permissionDenied e [agr] s action a)  \/
(queryInconsistent e [agr] s action a) \/
(permissionUnregulated e [agr] s action a).
```

We will then augment ODRL0 with the constructs we omitted from the full ODRL (resulting in what we have earlier called ODRL1 or ODRL2) including the troublesome construct *not[policySet]* and attempt to prove that the decidability results remain intact. There is a chance that a proof is not possible due to particulars of the Coq encoding we have used, in in which case, we will adjust our encoding.

9.3   SELinux

We started out by looking at DRELs and specifically ODRL where rights expressions are used to arbitrate access to assets under conditions. Recall that the main construct in ODRL is the agreement which specifies users, asset(s) and policies (as part of policy sets) whereby controls on users' access to the assets are described. This is reminiscent of how access control conditions are expressed in access control policy languages such as XACML and SELinux.

While XACML is a high-level and platform independent access control system SELinux is platform dependent (e.g. Linux based) and low-level. SELinux enhances the Discretionary Access Control (DAC) that most unix based systems employ by Mandatory Access Control (MAC) where designed access control policies are applied throughout the system possibly overriding whatever DAC is in place by the system users.

SELinux uses Linux's extended file attributes to attach a *security context* to passive entities (e.g. files, directories, sockets) and also to each active entity typically a Linux user space process. Security context is a data structure that is composed of a user, a role and a domain (or type). While users can map directly to ordinary user names they can also be defined separately. Roles are meant to group users and add flexibility when assigning permissions and are the basis for RBAC. Finally domains or types are the basis for defining common access control requirements for both passive and active entities.

The enforcement of SELinux policies are performed by the *security server*. Whenever a security operation is requested from user space by a system call, the security server is invoked to arbitrate the operation and either allow the operation or to deny it. Each operation is identified by two pieces of information: an object class (e.g. file) and a

permission (e.g. read, write). When an operation is requested to be performed on an object, the class and the permissions associated with the object along with security contexts of the source (typically the source entity is a process) and the object are passed to the security server. The security server consults the loaded policy (loaded at boot time) and allows or denies the access request [SSS04].

9.4   SELinux Policy Language

The SELinux policy has four different kinds of statements: declarations, rules, constraints and assertions [ALP03]. Assertions are compile time checks that the *checkpolicy* tool performs at compile time. The other three kinds of statements however are evaluated at run-time.

Declaration statements are used to declare a type, a role and a type. First types are declared using a type declaration statement. Roles are declared and authorized for particular domains (types) through role declarations, and finally user declarations are used to define each user and to specify the set of authorized roles for each of these users (see 9.2). Note that in the listings below we will present a simplified and modified version of the official SELinux syntax in order to (re-)use the ODRL Coq framework.

Listing 9.2: Declarations

```
'type' T ';'                                    ; type T

'role' R T ';'                                  ; role R is associated with type T

'user' U R ';'                                  ; user U is associated with role R
```

Rule statements define access vector rules. Access vector (AV) rules (see listing 9.3) specify which operations are allowed and whether to audit (log). Any operation not covered by AV rules are denied and all denied operations are logged. The semantics of the AV rule with avkind *allow* is: processes with type *T1* are allowed to perform operations in *P* on objects with class *C* and type *T2* (with avkind=*deny* meaning not to allow).

Listing 9.3: AV Rule

```
<avRule> → <avkind> T1 T2:C P ';'

<avkind> → 'allow' | 'deny'
```

When a process changes security context, the role may change, assuming a *role transition* exists relating the old and the new roles. There is a related AV rule called the *type transition* rule where a process with type T1 is allowed/denied to transition to type T2 when C=process and P=transition (see 9.4).

Listing 9.4: Type Transition and Role-Allow Rules

```
<avkind> T1 T2:process transition ';'


<role_transition_rule> → <avkind> R1 R2 ';'
   ; when a process changes security contexts this rule must hold
```

Constraints are additional conditions on permissions in the form of boolean expressions that must hold in order for the specified permissions to be allowed (see listing 9.5. Whenever a permission is requested on an object class C, the security server checks that the constraints hold.

Listing 9.5: Constrain Rule

```
<constraint> → 'constrain' C, P, <expr> ';'


<expr> → 'not' <expr> | <expr> 'and' <expr> | <expr> or <expr> | U1 <
    op> U2 | T1 <op> T2 | R1 <op> R2


<op> → '==' | '!='
```

9.5 Agreements in SELinux

As with ODRL we will start by limiting the policy language to only allow AV rules. As mentioned earlier an operation not covered by a allow rule is denied by SELinux. We will make up explicit *deny* rules therefore an agreement is defined to be a combination of allow and deny rules. Allow and deny rules as mappings are defined in listing 9.6.

Listing 9.6: 'allow'/'deny' Rule as a Mapping

```
AV rule : $T \times (T \times C) \to 2^P$
```

Listing 9.7: SELinux Agreement

```
<agreement> ::= <avRule> ';' <agreement>
```

9.6 Environments

Environments are collections of *role-type* and *user-role* relations. A role-type relation $role(R, T)$ simply associates a role with a type. A user-role relation $user(U, R)$ associates a user with a role. An environment is consistent with respect to a security context $< T, R, U >$, if and only if $role(R, T)$ and $user(U, R)$ relations hold in the environment.

9.7 Queries in SELinux

The decision problem in SELinux access control is whether an entity with security context $<T1, R1, U1>$ may perform action $P1$ to entity with object class $C1$ with security context $<T2, R2, U2>$.

To answer such queries we use the authorization relation *auth(C, P, T1, R1, U1, T2, R2, U2)* which is equivalent to the *Permitted* predicate we saw earlier for ODRL.

Listing 9.8: f_q^+ for SELinux

$allow(T1, T2, C, P) \land (E \; consistent \; wrt \; \texttt{<T1, R1, U1>} \land \texttt{<T2, R2, U2>}) \land$
$\quad (((C, P) == (process, transition)) \implies allow(R1, R2))$
$\quad \implies auth(C, P, T1, R1, U1, T2, R2, U2)$

Listing 9.9: f_q^- for SELinux

$deny(T1, T2, C, P) \lor \neg(E \; consistent \; wrt \; \texttt{<T1, R1, U1>} \land \texttt{<T2, R2, U2>}) \lor$
$\quad (((C, P) == (process, transition)) \implies deny(R1, R2))$
$\quad \implies \neg auth(C, P, T1, R1, U1, T2, R2, U2)$

9.8 Decidability of Queries in SELinux

In this thesis we will be investigating the question of decidability for answering queries in SELinux policies based on the same four outcomes we encountered earlier in 5.3 namely error, permitted, denied and "not applicable". We will first state a decidability theorem similar to the theorem in listing 9.1 (minor adjustments may be needed to allow for differences with SELinux policy language) and present a proof for it in Coq. The literature in the SELinux implies only two outcomes are possible: permitted or denied. We will next attempt to prove this conjecture in Coq. Finally we will add constraint relations (see listing 9.10 and listing 9.11) to SELinux policies (which we have not included so far) and prove the same decidability results again for the augmented policy.

Listing 9.10: f_q^+ for SELinux

$allow(T1, T2, C, P) \land constrain(C, P, T1, R1, U1, T2, R2, U2) \land (E \; consistent \; wrt \; \texttt{<}$
$\quad \texttt{T1, R1, U1>} \land \texttt{<T2, R2, U2>}) \land (((C, P) == (process, transition)) \implies$
$\quad allow(R1, R2)) \implies auth(C, P, T1, R1, U1, T2, R2, U2)$

Listing 9.11: f_q^- for SELinux

$deny(T1, T2, C, P) \lor \neg \; constrain(C, P, T1, R1, U1, T2, R2, U2) \lor \neg(E \; consistent \; wrt$
$\quad \texttt{<T1, R1, U1>} \land \texttt{<T2, R2, U2>}) \lor (((C, P) == (process, transition)) \implies$
$\quad deny(R1, R2)) \implies \neg auth(C, P, T1, R1, U1, T2, R2, U2)$

References

[ALP03]   Myla Archer, Elizabeth I. Leonard, and Matteo Pradella. Analyzing security-enhanced linux policy specifications. In *4th IEEE International Workshop on Policies for Distributed Systems and Networks (POLICY 2003), 4-6 June 2003, Lake Como, Italy*, page 158. IEEE Computer Society, 2003.

[ANP+03]   Anne Anderson, Anthony Nadalin, B Parducci, D Engovatov, H Lockhart, M Kudo, P Humenn, S Godik, S Anderson, S Crocker, et al. eXtensible access control markup language (xacml) version 1.0. *OASIS*, 2003.

[ASLZ08]   Masoom Alam, Jean-Pierre Seifert, Qi Li, and Xinwen Zhang. Usage control platformization via trustworthy SELinux. In *Proceedings of the 2008 ACM symposium on Information, computer and communications security*, pages 245–248. ACM, 2008.

[BC04]   Yves Bertot and Pierre Castéran. *Interactive Theorem Proving and Program Development. Coq'Art: The Calculus of Inductive Constructions*. Springer Science & Business Media, 2004. http://www.labri.fr/perso/casteran/CoqArt/index.html.

[Ber10]   Yves Bertot. Coq in a hurry. Technical report, May 2010. http://cel.archives-ouvertes.fr/inria-00001173/PDF/coq-hurry.pdf.

[CPR04]   Geoff Collier, Harry Piccariello, and Robby Robson. A digital rights management ecosystem model for the education community. *ECAR (forthcoming), May*, 2004.

[dt04]   The Coq development team. *The Coq proof assistant reference manual*. LogiCal Project, 2004. Version 8.0.

[HKS04]   Markus Holzer, Stefan Katzenbeisser, and Christian Schallhart. Towards formal semantics for ODRL. In Renato Iannella and Susanne Guth, editors, *Proceedings of the First International Workshop on the Open Digital Rights Language (ODRL), Vienna, Austria, April 22-23, 2004*, pages 137–148, 2004.

[Hue92]   Gérard P. Huet. The gallina specification language: A case study. In R. K. Shyamasundar, editor, *Foundations of Software Technology and Theoretical Computer Science, 12th Conference, New Delhi, India, December 18-20, 1992,*

Proceedings, volume 652 of *Lecture Notes in Computer Science*, pages 229–240. Springer, 1992.

[HW08]   Joseph Y. Halpern and Vicky Weissman. Using First-Order Logic to Reason about Policies. *ACM Transactions on Information and System Security*, 11(4), 2008.

[Ian02]   R. Iannella. Open digital rights language (ODRL) version 1.1. 2002. [Online; accessed 09-February-2015].

[JHM06]   Pramod A. Jamkhedkar, Gregory L. Heileman, and Iván Martínez-Ortiz. The problem with rights expression languages. In Moti Yung, Kaoru Kurosawa, and Reihaneh Safavi-Naini, editors, *Proceedings of the Sixth ACM Workshop on Digital Rights Management, Alexandria, VA, USA, October 30, 2006*, pages 59–68. ACM, 2006.

[KG10]   Andreas Kasten and RÃijdiger Grimm. Making the semantics of ODRL and URM explicit using web ontologies. In *The 8th International Workshop for Technical, Economic and Legal Aspects of Business Models for Virtual Goods*, pages 77–91, Namur, Belgium, 2010.

[MPT12]   Massimiliano Masi, Rosario Pugliese, and Francesco Tiezzi. Formalisation and implementation of the XACML access control mechanism. In Gilles Barthe, Benjamin Livshits, and Riccardo Scandariato, editors, *Engineering Secure Software and Systems - 4th International Symposium, ESSoS 2012, Eindhoven, The Netherlands, February, 16-17, 2012. Proceedings*, volume 7159 of *Lecture Notes in Computer Science*, pages 60–74. Springer, 2012.

[MRB98]   Robert Martin, Dirk Riehle, and Frank Buschmann. *Pattern Languages of Program Design 3*, volume 256. Addison-Wesley, 1998.

[MRD09]   Xavier Maroñas, Eva Rodríguez, and Jaime Delgado. An architecture for the interoperability between rights expression languages based on XACML. In *Proceedings of the 7th International Workshop for technical, economic and legal aspects of business models for virtual goods, Virtual Goods*, 2009.

[NIS09]   NIST.   A SURVEY OF ACCESS CONTROL MODELS.   http://csrc.nist.gov/news_events/privilege-management-workshop/PvM-Model-Survey-Aug26-2009.pdf, 2009.   [Online; accessed 16-August-2015].

[PCG+11]   Benjamin C. Pierce, Chris Casinghino, Michael Greenberg, Vilhelm Sjöberg, and Brent Yorgey. *Software Foundations*. University of Pennsylvania, 2011. http://www.cis.upenn.edu/~bcpierce/sf/.

[Pie02]   Benjamin C. Pierce. *Types and Programming Languages*. MIT Press, Cambridge, MA, USA, 2002.

[PRD05]   Jose Prados, Eva Rodriguez, and Jaime Delgado. Interoperability between different rights expression languages and protection mechanisms. In *Automated Production of Cross Media Content for Multi-Channel Distribution, 2005. AXMEDIS 2005. First International Conference on*, pages 8–pp. IEEE, 2005.

[PW06]   Riccardo Pucella and Vicky Weissman. A formal foundation for ODRL. *CoRR*, abs/cs/0601085, 2006.

[Sis15]   Bahman Sistany. Source code for accpl. https://github.com/bsistany/yacpl.git, 2015.

[SM12]   Michel St-Martin. *A verified algorithm for detecting conflicts in XACML access control rules.* PhD thesis, Carleton University, 2012.

[SMF13]   Michel St-Martin and Amy Felty. A verified algorithm for detecting conflicts in XACML access control rules. In *Informal Proceedings of the First International Workshop on Automated Reasoning in Security and Software*, pages 4–11, 2013. http://www.site.uottawa.ca/~afelty/dist/arsec13.pdf.

[SS09]   Nicholas Paul Sheppard and Reihaneh Safavi-Naini. On the operational semantics of rights expression languages. In Ehab Al-Shaer, Hongxia Jin, and Gregory L. Heileman, editors, *Proceedings of the 9th ACM Workshop on Digital Rights Management, Chicago, Illinois, USA, November 9, 2009*, pages 17–28. ACM, 2009.

[SSS04]   Beata Sarna-Starosta and Scott D. Stoller. Policy analysis for security-enhanced linux. In *Proceedings of the 2004 Workshop on Issues in the Theory of Security (WITS)*, pages 1–12. Citeseer, 2004.

[SVLS02]   Stephen Smalley, Chris Vance, NAI Labs, and Wayne Salamon. Implementing SELinux as a Linux security module. In *USENIX Technical Conference*, 2002. [Online; accessed 16-February-2015].

[TK06]   Michael Carl Tschantz and Shriram Krishnamurthi. Towards reasonability properties for access-control policy languages. In David F. Ferraiolo and Indrakshi Ray, editors, *SACMAT 2006,11th ACM Symposium on Access Control Models and Technologies, Lake Tahoe, California, USA, June 7-9, 2006, Proceedings*, pages 160–169. ACM, 2006.

[Wik15a]   Wikipedia. FairPlay — Wikipedia, The Free Encyclopedia. http://en.wikipedia.org/wiki/FairPlay, 2015. [Online; accessed 16-February-2015].

[Wik15b]   Wikipedia. Open Web Platform — Wikipedia, The Free Encyclopedia. http://en.wikipedia.org/wiki/Open_Web_Platform, 2015. [Online; accessed 16-February-2015].

[Wik15c]   Wikipedia. PlayReady — Wikipedia, The Free Encyclopedia. http://en.wikipedia.org/wiki/PlayReady, 2015. [Online; accessed 16-February-2015].

[WLD+02] Xin Wang, Guillermo Lao, Thomas DeMartini, Hari Reddy, Mai Nguyen, and Edgar Valenzuela. XrML - eXtensible rights markup language. In Michiharu Kudo, editor, *Proceedings of the 2002 ACM Workshop on XML Security, Fairfax, VA, USA, November 22, 2002*, pages 71–79. ACM, 2002.